

Operating Systems

Concurrency and Thread API (Chapter 26 ~ 27)

Young-Woo Kwon


CONCURRENCY

Concurrent/Parallel Programs

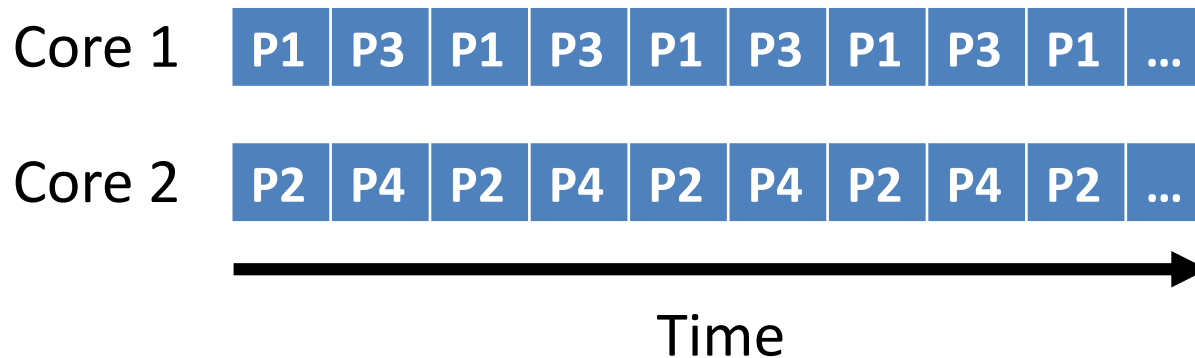
- To execute these programs we need to
 - **Create several processes** that execute in parallel (or concurrently)
 - Map each process to the **same address space to share data**
 - They are all part of the same computation
 - Have the **OS schedule** these processes in parallel
- This situation is very inefficient. Why?

Concurrency vs. Parallelism

- Concurrent execution on a single-core system:



- Parallel execution on a dual-core system:

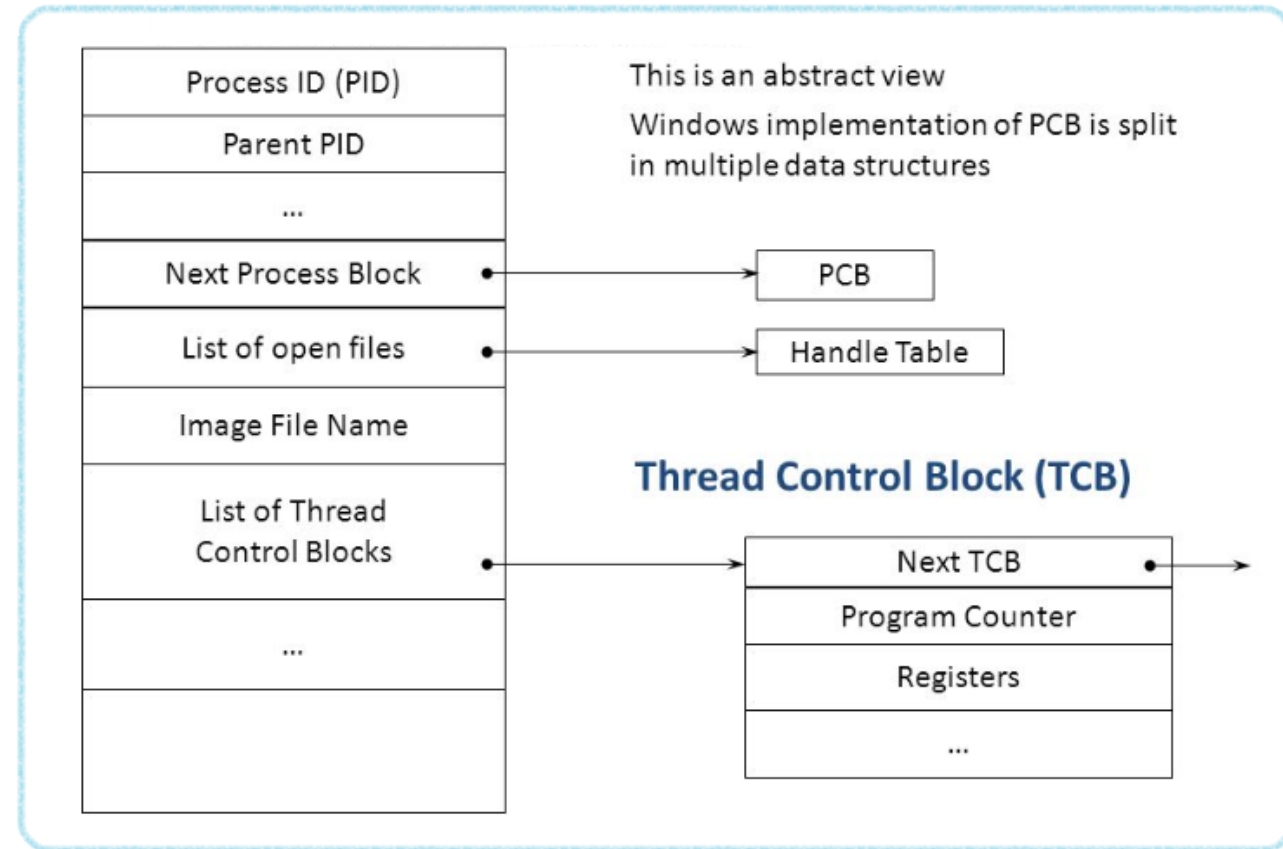


Thread

- A **thread** is a single execution sequence that represents a separately schedulable task
- Multi-threaded program
 - A multi-threaded program has **more than one point of execution.**
 - Multiple PCs (Program Counter)
 - They **share** the share the same **address space.**

Thread Control Block (TCB)

- Thread ID
- Thread state
- Pointer to parent PCB
- PC/registers for thread
- Stack location in memory map.

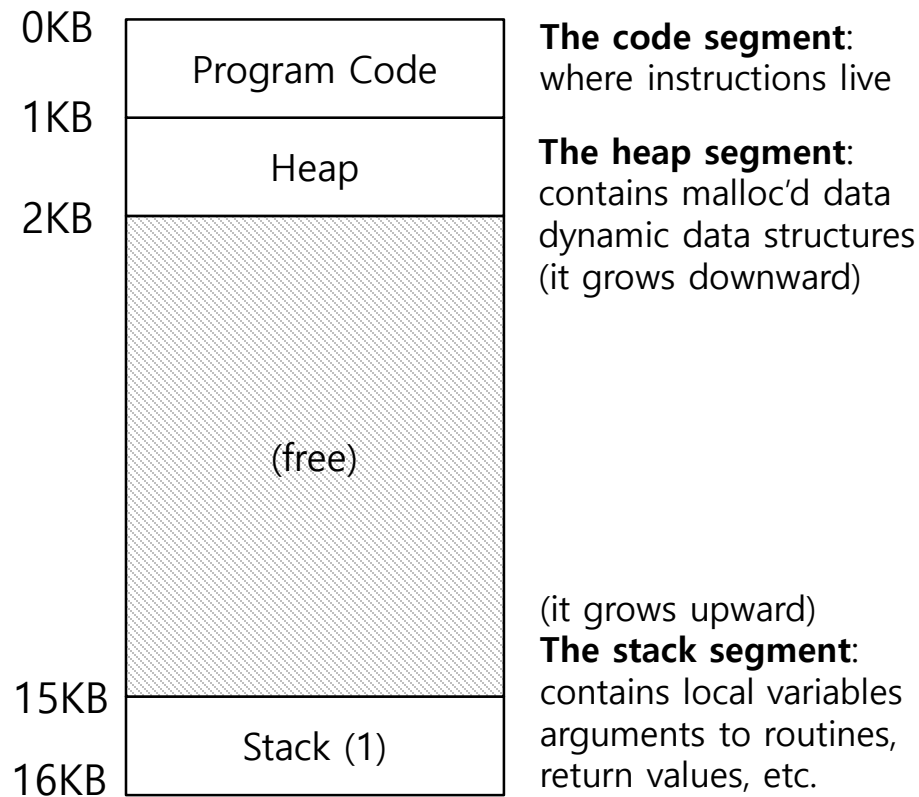


Context switch between threads


- Each thread has its own program counter and set of registers.
 - One or more **thread control blocks(TCBs)** are needed to store the state of each thread.
- When switching from running one (T1) to running the other (T2),
 - The register state of T1 be saved.
 - The register state of T2 restored.
 - The **address space remains** the same.

The stack of the relevant thread

- There will be **one stack per thread**.

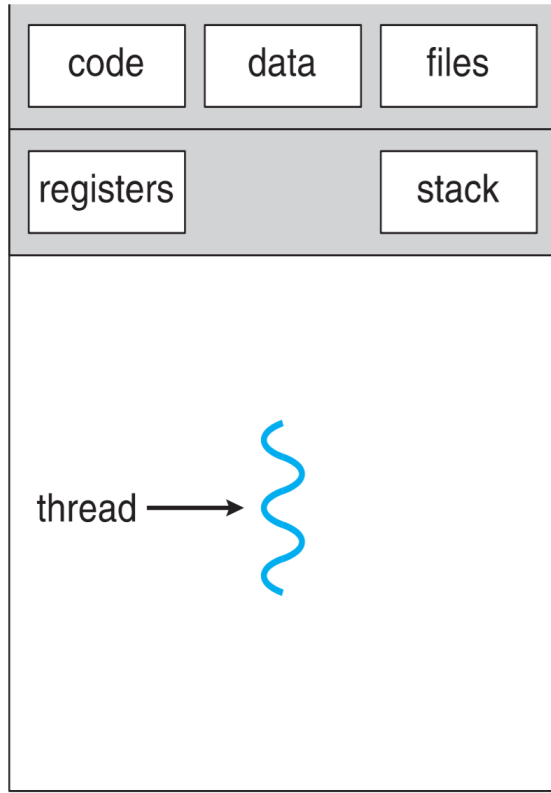


**A Single-Threaded
Address Space**




**Two threaded
Address Space**

Single vs. Multithreaded Processes



single-threaded process



multithreaded process

Why Use Threads?

- Parallelism
 - One thread per CPU can make programs **run faster** on multiple processors
- I/O overlapping
 - **Avoid blocking program progress** due to slow I/O
 - While one thread in your program waits (i.e., blocked waiting for I/O), the CPU scheduler can switch to other threads, which are ready to run and do something useful.
 - Similar to the effect of multiprocessing

Why Use Threads?

- You could use multiple processes instead of threads.
 - Processes are a more sound choice for logically separate tasks
- But,
 - **Process creation is heavy-weight** while thread creation is lightweight
 - **Threads share an address space** and thus make it easy to share data
 - Can simplify code, **increase efficiency**

Threads are “lighter weight” than processes

- To make a new thread, we only need a **stack**.
 - Can share all pre-existing resources.
 - Done at user-level without system calls
- To make a new process, we have to talk to the **operating system**
 - Make a new address space
 - Inherit resources from the original space
 - Compete for the same underlying global scheduler
 - Can easily run out, or thrash

Benefits

- **Responsiveness**
 - may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing**
 - threads share resources of process, easier than shared memory or message passing
- **Economy**
 - cheaper than process creation, thread switching lower overhead than context switching
- **Scalability**
 - Threads can take advantage of multiprocessor architectures

An Example: Thread Creation

```
1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>

4  void *mythread (void *arg) {
5      printf ("%s\n", (char *) arg);
6      return NULL;
7  }
8
9  int main (int argc, char *argv[]) {
10     pthread_t p1, p2;
11     int rc;
12     printf("main: begin\n");
13     pthread_create(&p1, NULL, mythread, "A");
14     pthread_create(&p2, NULL, mythread, "B");
15     // join waits for three threads to finish
16     pthread_join(p1, NULL);
17     pthread_join(p2, NULL);
18     printf("main: end\n");
19     return 0;
20 }
```

Thread Trace (1)

main	Thread 1	Thread 2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
waits for T1		
<hr/>		
	runs	
	prints "A"	
	returns	
<hr/>		
waits for T2		
<hr/>		
		runs
		prints "B"
		returns
<hr/>		
prints "main: end"		

Thread Trace (2)

main	Thread 1	Thread 2
starts running		
prints "main: begin"		
creates Thread 1		
<hr/>		
	runs	
	prints "A"	
	returns	
<hr/>		
creates Thread 2		
<hr/>		
		runs
		prints "B"
		returns
<hr/>		
waits for T1		
returns immediately; T1 is done		
waits for T2		
returns immediately; T2 is done		
prints "main: end"		

Thread Trace (3)

main	Thread 1	Thread 2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
<hr/>		
		runs
		prints "B"
		returns
<hr/>		
waits for T1		
<hr/>		
	runs	
	prints "A"	
	returns	
<hr/>		
waits for T2		
returns immediately; T2 is done		
prints "main: end"		

Issues when using threads

```
1  #include <stdio.h>
2  #include <assert.h>
3  #include <pthread.h>jjj
4  int counter = 0;
5  void *mythread (void *arg) {
6      printf ("%s\n", (char *) arg);
7      for(int i = 0; i < 1e7; i++)
8          counter = counter + 1;
9      printf ("%s: done\n", (char *) arg);
10     return NULL;
11 }
12
13 int main (int argc, char *argv[]) {
14     pthread_t p1, p2;
15     int rc;
16     printf("main: begin\n");
17     pthread_create(&p1, NULL, mythread, "A");
18     pthread_create(&p2, NULL, mythread, "B");
19     // join waits for thre threads to finish
20     pthread_join(p1, NULL);
21     pthread_join(p2, NULL);
22     printf("main: end\n");
23     return 0;
24 }
```

Result

```
main: begin (counter = 0)
```

```
A
```

```
B
```

```
B: done
```

```
A: done
```

```
main: end (counter = 10313459)
```

```
main: begin (counter = 0)
```

```
A
```

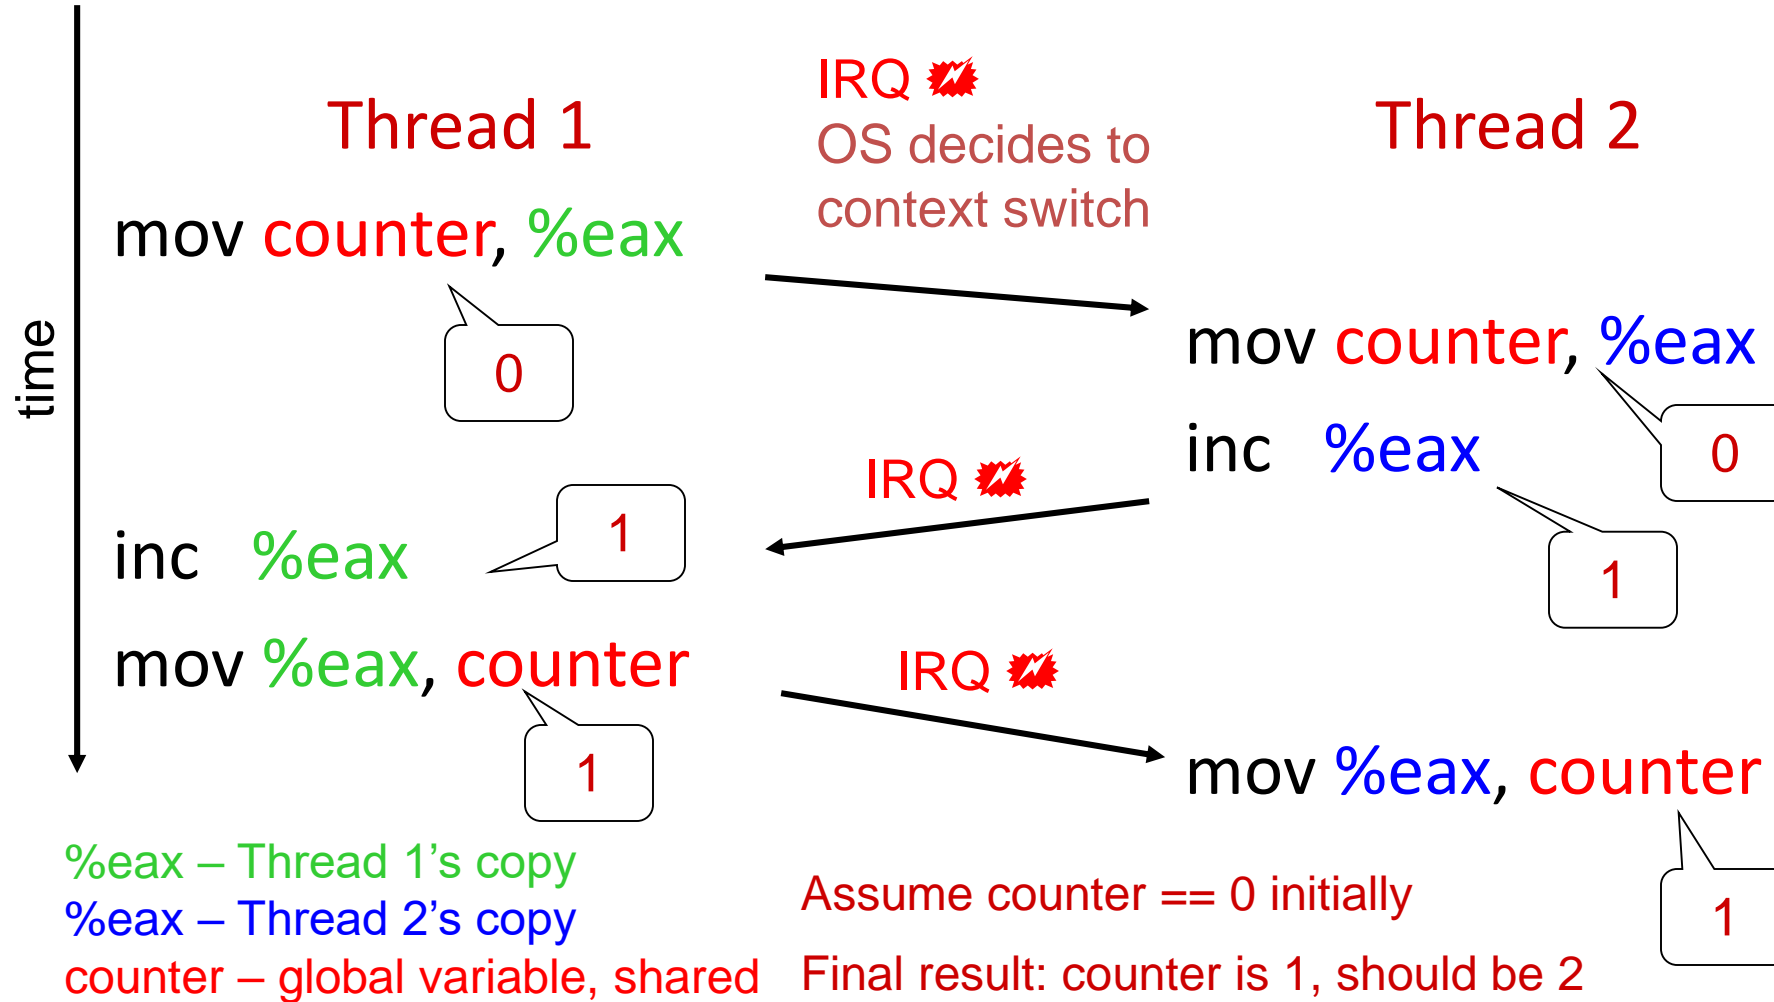
```
B
```

```
A: done
```

```
B: done
```

```
main: end (counter = 10201521)
```

Race Conditions



Race Conditions



- Race condition
 - Two threads “race” to execute code and update shared (dependent) data
 - Errors emerge based on the ordering of operations, and the scheduling of threads
 - Thus, errors are nondeterministic

Critical Sections

- Classical definition of a critical section:
 - “A piece of code that **accesses a shared resource** that must not be concurrently accessed by more than one thread of execution.”
 - Multiple threads executing critical section can result in a race condition.
 - Need to support **atomicity** for critical sections (**mutual exclusion**)
- Two problems
 - Code was not designed for concurrency
 - Shared resource (data) does not support concurrent access

Wish for Atomicity


- We need more powerful instructions
 - do exactly whatever we needed done in a single step → atomic
 - remove the possibility of an untimely interrupt



Mutexes for Atomicity

- Mutual exclusion lock (mutex) is a construct that can enforce atomicity in code


```
m = mutex_create();  
...  
mutex_lock(m);  
// do some stuff  
mutex_unlock(m);
```



One More Problem: Waiting For Another

- Another common interaction
 - One thread must wait for another to complete some action before it continues
 - e.g., when a process performs a disk I/O and is put to sleep; when the I/O completes, the process needs to be roused from its slumber so it can continue

Thread Lifecycle



27. Thread API

Thread Creation

- How to create and control threads?

```
#include <pthread.h>

int
pthread_create(      pthread_t*      thread,
                    const pthread_attr_t* attr,
                    void*             (*start_routine) (void*),
                    void*             arg);
```

- `thread`: Used to interact with this thread.
- `attr`: Used to specify any attributes this thread might have.
 - Stack size, Scheduling priority, ...
- `start_routine`: the function this thread start running in.
- `arg`: the argument to be passed to the function (`start routine`)
 - *a void pointer* allows us to pass in *any type of* argument.

Thread Creation (Cont.)

- If `start_routine` instead requires another type argument, the declaration would look like this:
 - An integer argument:

```
int  
pthread_create(..., // first two args are the same  
                void*   (*start_routine)(int),  
                int     arg);
```

- Return an integer:

```
int  
pthread_create(..., // first two args are the same  
                int     (*start_routine)(void*),  
                void*   arg);
```

Example: Creating a Thread

```
#include <pthread.h>

typedef struct __myarg_t {
    int a;
    int b;
} myarg_t;

void *mythread(void *arg) {
    myarg_t *m = (myarg_t *) arg;
    printf("%d %d\n", m->a, m->b);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p;
    int rc;

    myarg_t args;
    args.a = 10;
    args.b = 20;
    rc = pthread_create(&p, NULL, mythread, &args);
    ...
}
```

Wait for a thread to complete

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- thread: Specify which thread *to wait for*.
- value_ptr: A pointer to the return value
 - Because `pthread_join()` routine changes the value, you need to **pass the pointer** to that value.

Example: Waiting for Thread Completion

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include <assert.h>
4  #include <stdlib.h>
5
6  typedef struct __myarg_t {
7      int a;
8      int b;
9  } myarg_t;
10
11 typedef struct __myret_t {
12     int x;
13     int y;
14 } myret_t;
15
16 void *mythread(void *arg) {
17     myarg_t *m = (myarg_t *) arg;
18     printf("%d %d\n", m->a, m->b);
19     myret_t *r = malloc(sizeof(myret_t));
20     r->x = 1;
21     r->y = 2;
22     return (void *) r;
23 }
24
```


Example: Waiting for Thread Completion (Cont.)

```
25  int main(int argc, char *argv[]) {
26      int rc;
27      pthread_t p;
28      myret_t *m;
29
30      myarg_t args;
31      args.a = 10;
32      args.b = 20;
33      pthread_create(&p, NULL, mythread, &args);
34      pthread_join(p, (void **) &m); // this thread has been
                                     // waiting inside of the
                                     // pthread_join() routine.
35      printf("returned %d %d\n", m->x, m->y);
36      return 0;
37 }
```

Example: Dangerous code

- Be careful with **how values are returned** from a thread.

```
1  void *mythread(void *arg) {  
2      myarg_t *m = (myarg_t *) arg;  
3      printf("%d %d\n", m->a, m->b);  
4      myret_t r;  
5      r.x = 1;  
6      r.y = 2;  
7      return (void *) &r;  
8  }
```

- The variable `r` is allocated in a stack space. So?

Example: Simpler Argument Passing to a Thread

- Just passing in a single value

```
1  void *mythread(void *arg) {
2      int m = (int) arg;
3      printf("%d\n", m);
4      return (void *) (arg + 1);
5  }
6
7  int main(int argc, char *argv[]) {
8      pthread_t p;
9      int rc, m;
10     pthread_create(&p, NULL, mythread, (void *) 100);
11     pthread_join(p, (void **) &m);
12     printf("returned %d\n", m);
13     return 0;
14 }
```

Locks

- Provide **mutual exclusion** to a **critical section**

- Interface

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- Usage (*w/o lock initialization and error check*)

```
pthread_mutex_t lock;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

- No other thread holds the lock → the thread will acquire the lock and **enter the critical section**.
- If another thread hold the lock → the thread will **not return from the call** until it has acquired the lock.

Condition Variables

- **Condition variables** are useful when some kind of **signaling** must take place between threads.

```
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex);  
int pthread_cond_signal(pthread_cond_t *cond);
```

– pthread_cond_wait:

- Put the calling thread to sleep.
- Wait for some other thread to signal it.

– pthread_cond_signal:

- Unblock at least one of the threads that are blocked on the condition variable

Condition Variables (Cont.)

- The waiting thread **re-checks** the condition **in a while loop**, instead of a simple if statement.

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
pthread_cond_t init = PTHREAD_COND_INITIALIZER;  
  
pthread_mutex_lock(&lock);  
while (initialized == 0)  
    pthread_cond_wait(&init, &lock);  
pthread_mutex_unlock(&lock);
```

- Without rechecking, the waiting thread will continue thinking that the condition has changed *even though it has not*.

Condition Variables (Cont.)

- Don't ever to this.x
 - A thread calling wait routine:

```
while(initialized == 0)  
    ; // spin
```

- A thread calling signal routine:

```
initialized = 1;
```

- It performs poorly in many cases. → just wastes CPU cycles.
- It is error prone.

Implementing a Lock

```
do {  
    while (flag);  
    flag = true;  
    critical section  
    flag = false;  
    remainder section  
} while (true);
```


Problem?

```
flag = false;
```

```
do {  
    while (flag);  
    flag = true;  
        critical section  
    flag = false;  
        remainder section  
} while (true);
```

```
do {  
    while (flag);  
    flag = true;  
        critical section  
    flag = false;  
        remainder section  
} while (true);
```