

Operating Systems

Lock (Chapter 28 ~ 29)

Dr. Young-Woo Kwon

Definitions

- **Race condition:** output of a concurrent program depends on the order of operations between threads
- **Mutual exclusion:** only one thread does a particular thing at a time
- **Critical section:** piece of code that only one thread can execute at once
- **Lock:** prevent someone from doing something

Locks: The Basic Idea

- Ensure that any **critical section** executes as if it were a **single atomic instruction**.
 - An example: the canonical update of a shared variable

```
balance = balance + 1;
```

- Add some code around the critical section

```
1  lock_t mutex; // some globally-allocated lock 'mutex'
2  ...
3  lock(&mutex);
4  balance = balance + 1;
5  unlock(&mutex);
```

Locks: The Basic Idea

- Lock variable holds the state of the lock.
 - **available** (or **unlocked** or **free**)
 - No thread holds the lock.
 - **acquired** (or **locked** or **held**)
 - Exactly one thread holds the lock and presumably is in a critical section.

The semantics of the lock()

- `lock()`
 - **Try to** acquire the lock.
 - If no other thread holds the lock, the thread will **acquire** the lock.
 - **Enter** the *critical section*.
 - This thread is said to be the owner of the lock.
 - Other threads are *prevented from* entering the critical section while the first thread that holds the lock is in there.

Pthread Locks - mutex

- The name that the POSIX library uses for a lock.
 - Used to provide **mutual exclusion** between threads.

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
2  
3  Pthread_mutex_lock(&lock); // wrapper for pthread_mutex_lock()  
4  balance = balance + 1;  
5  Pthread_mutex_unlock(&lock);
```

- We may be using *different locks* to protect *different variables* → Increase **concurrency** (a more **fine-grained** approach).

Building A Lock

- Efficient locks provided mutual exclusion at **low cost**.
- Building a lock need some help from the **hardware** and the **OS**.

Evaluating locks – Basic criteria

- **Mutual exclusion**

- Does the lock work, preventing multiple threads from entering a *critical section*?

- **Fairness**

- Does each thread contending for the lock get a fair shot at acquiring it once it is free? (Starvation)

- **Performance**

- The time overheads added by using the lock

Why hardware support needed?

- **First attempt:** Using a *flag* denoting whether the lock is held or not.
 - The code below has problems.

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 → lock is available, 1 → held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10         ; // spin-wait (do nothing)
11     mutex->flag = 1; // now SET it !
12 }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

Why hardware support needed? (Cont.)

- **Problem 1: No Mutual Exclusion** (assume `flag=0` to begin)

Thread1	Thread2
<pre>call lock() while (flag == 1) interrupt: switch to Thread 2</pre>	<pre>call lock() while (flag == 1) flag = 1; interrupt: switch to Thread 1</pre>
<pre>flag = 1; // set flag to 1 (too!)</pre>	

- **Problem 2: Spin-waiting** wastes time waiting for another thread.
- So, we need an **atomic instruction** supported by **Hardware!**
 - *test-and-set* instruction, also known as *atomic exchange*

Test And Set (Atomic Exchange)

- An instruction to support the creation of simple locks

```
1  int TestAndSet(int *ptr, int new) {  
2      int old = *ptr;  // fetch old value at ptr  
3      *ptr = new;      // store 'new' into ptr  
4      return old;      // return the old value  
5  }
```

- **return**(testing) old value pointed to by the `ptr`.
- *Simultaneously* **update**(setting) said value to `new`.
- This sequence of operations is **performed atomically**.

A Simple Spin Lock using test-and-set

```
1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0 indicates that lock is available,
7      // 1 that it is held
8      lock->flag = 0;
9  }
10
11 void lock(lock_t *lock) {
12     while (TestAndSet(&lock->flag, 1) == 1)
13         ; // spin-wait
14 }
15
16 void unlock(lock_t *lock) {
17     lock->flag = 0;
18 }
```

- **Note:** To work correctly on a single CPU, it requires a preemptive scheduler.
 - **Why?**

Evaluating Spin Locks

- **Correctness:** *yes*
 - The spin lock only allows a single thread to entry the critical section.
- **Fairness:** *no*
 - Spin locks don't provide any fairness guarantees.
 - Indeed, a thread spinning may spin *forever*.
- **Performance:**
 - In the single CPU, performance overheads can be quire *painful*.
 - If **the number of threads roughly equals the number of CPUs**, spin locks work *reasonably well*.

Compare-And-Swap (SPARC)

- Test whether the value at the address(`ptr`) is equal to `expected`.
 - If so, **update** the memory location pointed to by `ptr` with the `new` value.
 - In either case, **return** the actual value at that memory location.

```
1  int CompareAndSwap(int *ptr, int expected, int new) {
2      int actual = *ptr;
3      if (actual == expected)
4          *ptr = new;
5      return actual;
6  }
```

Compare-and-Swap hardware atomic instruction (C-style)

```
1  void lock(lock_t *lock) {
2      while (CompareAndSwap(&lock->flag, 0, 1) == 1)
3          ; // spin
4  }
```

Spin lock with compare-and-swap

Compare-And-Exchange (x86)

- C-callable x86-version of compare-and-swap

```
1  char CompareAndSwap(int *ptr, int old, int new) {
2      unsigned char ret;
3
4      // Note that sete sets a 'byte' not the word
5      __asm__ __volatile__ (
6          " lock\n"
7          " cmpxchgl %2,%1\n"
8          " sete %0\n"
9          : "=q" (ret), "=m" (*ptr)
10         : "r" (new), "m" (*ptr), "a" (old)
11         : "memory");
12     return ret;
13 }
```

Fetch-And-Add

- **Atomically increment** a value while returning the old value at a particular address.

```
1  int FetchAndAdd(int *ptr) {  
2      int old = *ptr;  
3      *ptr = old + 1;  
4      return old;  
5  }
```

Fetch-And-Add Hardware atomic instruction (C-style)

Ticket Lock

- **Ticket lock** can be built with fetch-and add.
 - Ensure progress for all threads. → **fairness**

```
1  typedef struct __lock_t {
2      int ticket;
3      int turn;
4  } lock_t;
5
6  void lock_init(lock_t *lock) {
7      lock->ticket = 0;
8      lock->turn = 0;
9  }
10
11 void lock(lock_t *lock) {
12     int myturn = FetchAndAdd(&lock->ticket);
13     while (lock->turn != myturn)
14         ; // spin
15 }
16 void unlock(lock_t *lock) {
17     FetchAndAdd(&lock->turn);
18 }
```

So Much Spinning

- Hardware-based spin locks are **simple** and they work.
- In some cases, these solutions can be quite **inefficient**.
 - Any time a thread gets caught *spinning*, it **wastes an entire time slice** doing nothing but checking a value.

How To Avoid *Spinning*?
We'll need **OS Support** too!

A Simple Approach: Just Yield

- When you are going to spin, **give up the CPU** to another thread.
 - OS system call moves the caller from the *running state to the ready state*.
 - The cost of a **context switch** can be substantial and the **starvation** problem still exists.

```
1  void init() {  
2      flag = 0;  
3  }  
4  
5  void lock() {  
6      while (TestAndSet(&flag, 1) == 1)  
7          yield(); // give up the CPU  
8  }  
9  
10 void unlock() {  
11     flag = 0;  
12 }
```

Lock with Test-and-set and Yield

Using Queues: Sleeping Instead of Spinning

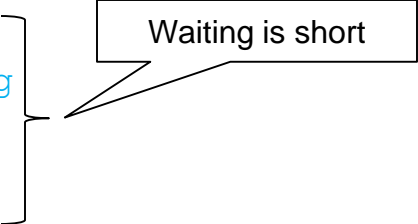
- **Queue** to keep track of which threads are waiting to enter the lock.
- `park()`
 - Put a calling thread to sleep
- `unpark(threadID)`
 - Wake a particular thread as designated by `threadID`.

Using Queues: Sleeping Instead of Spinning

```
typedef struct __lock_t {  
    int flag;           // lock is acquired or not  
    int guard;          // to protect the queue  
    queue_t *q;  
} lock_t;
```

Using Queues: Sleeping Instead of Spinning

```
1  typedef struct __lock_t { int flag; int guard; queue_t *q; } lock_t;
2
3  void lock_init(lock_t *m) {
4      m->flag = 0;
5      m->guard = 0;
6      queue_init(m->q);
7  }
8
9  void lock(lock_t *m) {
10     while (TestAndSet(&m->guard, 1) == 1)
11         ; // acquire guard lock by spinning
12     if (m->flag == 0) {
13         m->flag = 1; // lock is acquired
14         m->guard = 0;
15     } else {
16         queue_add(m->q, getpid());
17         m->guard = 0;
18         park();
19     }
20 }
21 ...
```



Lock With Queues, Test-and-set, Yield, And Wakeup

Using Queues: Sleeping Instead of Spinning

```
22 void unlock(lock_t *m) {  
23     while (TestAndSet(&m->guard, 1) == 1)  
24         ; // acquire guard lock by spinning  
25     if (queue_empty(m->q))  
26         m->flag = 0; // let go of lock; no one wants it  
27     else  
28         unpark(queue_remove(m->q)); // hold lock (for next thread!)  
29     m->guard = 0;  
30 }
```

Lock With Queues, Test-and-set, Yield, And Wakeup (Cont.)

Wakeup/waiting race

```
1  typedef struct __lock_t { int flag; int guard; queue_t *q; } lock_t;
2
3  void lock_init(lock_t *m) {
4      m->flag = 0;
5      m->guard = 0;
6      queue_init(m->q);
7  }
8
9  void lock(lock_t *m) {
10     while (TestAndSet(&m->guard, 1) == 1)
11         ; // acquire guard lock by spinning
12     if (m->flag == 0) {
13         m->flag = 1; // lock is acquired
14         m->guard = 0;
15     } else {
16         queue_add(m->q, getpid());
17         m->guard = 0;
18         park();
19     }
20 }
21 ...
```

What if a lock holder releases the lock just before the thread B executes “park()”?

- In case of releasing the lock (*thread A*) just before the call to `park()` (*thread B*) → Thread B would **sleep forever** (potentially).
- **Solaris** solves this problem by adding a third system call: `setpark()`.
 - By calling this routine, a thread can indicate it *is about to* park.
 - If it happens to be interrupted and another thread calls `unpark` before `park` is actually called, the subsequent `park` returns immediately instead of sleeping.

```
1      queue_add(m->q, getpid());
2      setpark(); // new code
3      m->guard = 0;
4      park();
```

Code modification inside of `lock()`


Spinning vs Blocking

- Some lock implementations combine spinning and blocking locks
- Blocking has a cost
 - Shouldn't block if lock becomes available in less time than it takes to block
- Strategy: spin for time it would take to block
 - Even in worst case, total cost for lock() is less than 2*block time

Two-Phase Locks


- A two-phase lock realizes that **spinning can be useful** if the lock *is about to* be released.
 - **First phase**
 - The lock spins for a while, *hoping that* it can acquire the lock.
 - If the lock is not acquired during the first spin phase, a second phase is entered,
 - **Second phase**
 - The caller is put to sleep.
 - The caller is only woken up when the lock becomes free later.

Critical Section Efficiency



$$Efficiency = \frac{T_c}{T_c + T_a + T_r}$$

- As processors get faster, CSE decreases because atomic instructions become relatively more expensive



[Source: McKenney, 2005](#)

29. Lock-based Concurrent Data Structures

Lock-based Concurrent Data structure

- Adding locks to a data structure makes the structure **thread safe**.
 - How locks are added determine both the **correctness** and **performance** of the data structure.

Example: Concurrent Counters without Locks

- Simple but not scalable

```
1      typedef struct __counter_t {  
2          int value;  
3      } counter_t;  
4  
5      void init(counter_t *c) {  
6          c->value = 0;  
7      }  
8  
9      void increment(counter_t *c) {  
10         c->value++;  
11     }  
12  
13     void decrement(counter_t *c) {  
14         c->value--;  
15     }  
16  
17     int get(counter_t *c) {  
18         return c->value;  
19     }
```

Example: Concurrent Counters with Locks

- Add a **single lock**.
 - The lock is acquired when calling a routine that manipulates the data structure.


```
1      typedef struct __counter_t {  
2          int value;  
3          pthread_lock_t lock;  
4      } counter_t;  
5  
6      void init(counter_t *c) {  
7          c->value = 0;  
8          Pthread_mutex_init(&c->lock, NULL);  
9      }  
10  
11     void increment(counter_t *c) {  
12         Pthread_mutex_lock(&c->lock);  
13         c->value++;  
14         Pthread_mutex_unlock(&c->lock);  
15     }  
16
```

Example: Concurrent Counters with Locks

```
(Cont.)
17     void decrement(counter_t *c) {
18         pthread_mutex_lock(&c->lock);
19         c->value--;
20         pthread_mutex_unlock(&c->lock);
21     }
22
23     int get(counter_t *c) {
24         pthread_mutex_lock(&c->lock);
25         int rc = c->value;
26         pthread_mutex_unlock(&c->lock);
27         return rc;
28     }
```


The performance costs of the simple approach

- Each thread updates a single shared counter.
 - Each thread updates the counter one million times.



Performance of
Traditional vs. Sloppy Counters

Synchronized counter scales poorly.

Perfect Scaling

- Even though more work is done, it is **done in parallel**.
- The time taken to complete the task is *not increased*.

Sloppy counter

- The sloppy counter works by representing ...
 - A single **logical counter** via numerous local physical counters, on per CPU core
 - A single **global counter**
 - There are **locks**:
 - One for each local counter and one for the global counter
- Example: on a machine with four CPUs
 - Four local counters
 - One global counter

The basic idea of sloppy counting

- When a thread running on a core wishes to increment the counter.
 - It increment its local counter.
 - Each CPU has its own local counter:
 - Threads across CPUs can update local counters *without contention*.
 - Thus counter updates are **scalable**.
 - The local values are periodically transferred to the global counter.
 - Acquire the global lock
 - Increment it by the local counter's value
 - The local counter is then reset to zero.

The basic idea of sloppy counting (Cont.)

- How often the local-to-global transfer occurs is determined by a threshold, S (sloppiness).
 - The smaller S :
 - The more the counter behaves like the *non-scalable counter*.
 - The bigger S :
 - The more scalable the counter.
 - The further off the global value might be from the *actual count*.

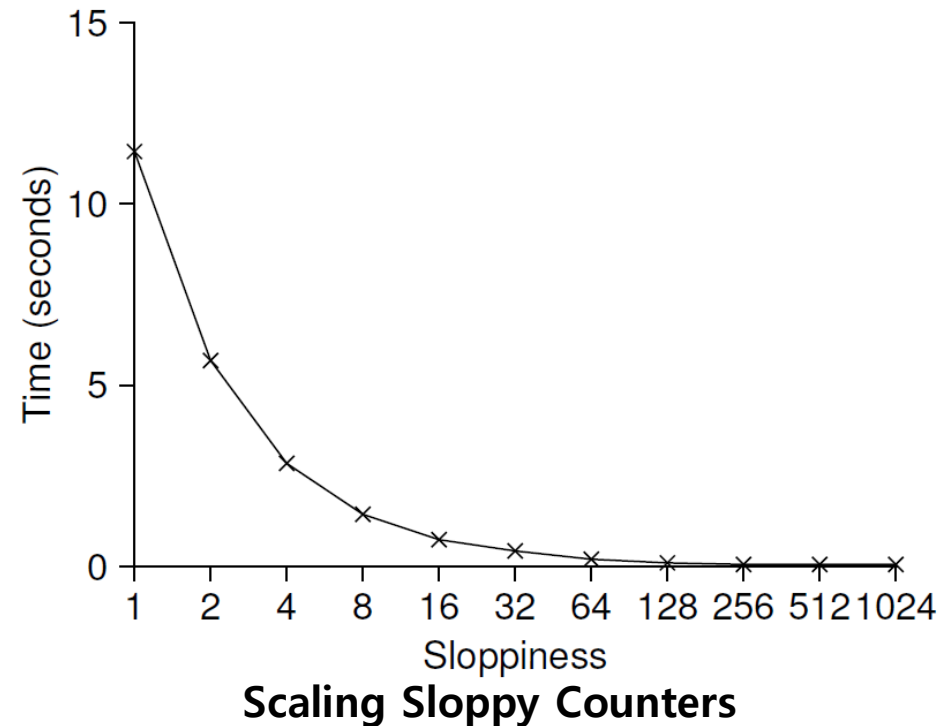
Sloppy counter example

- Tracing the Sloppy Counters
 - The threshold S is set to 5.
 - There are threads on each of four CPUs
 - Each thread updates their local counters $L_1 \dots L_4$.

Time	L1	L2	L3	L4	G
0	0	0	0	0	0
1	0	0	1	1	0
2	1	0	2	1	0
3	2	0	3	1	0
4	3	0	3	2	0
5	4	1	3	3	0
6	5 → 0	1	3	4	5 (from)
7	0	2	4	5 → 0	10 (from)

Importance of the threshold value S

- Each of four threads increments a counter 1 million times on four CPUs.
 - Low $S \rightarrow$ Performance is **poor**, The global count is always quite **accurate**.
 - High $S \rightarrow$ Performance is **excellent**, The global count **lags**.



Sloppy Counter Implementation

```
1     typedef struct __counter_t {
2         int global;           // global count
3         pthread_mutex_t glock; // global lock
4         int local[NUMCPUS];   // local count (per cpu)
5         pthread_mutex_t llock[NUMCPUS]; // ... and locks
6         int threshold;        // update frequency
7     } counter_t;
8
9     // init: record threshold, init locks, init values
10    //         of all local counts and global count
11    void init(counter_t *c, int threshold) {
12        c->threshold = threshold;
13
14        c->global = 0;
15        pthread_mutex_init(&c->glock, NULL);
16
17        int i;
18        for (i = 0; i < NUMCPUS; i++) {
19            c->local[i] = 0;
20            pthread_mutex_init(&c->llock[i], NULL);
21        }
22    }
23
```


Sloppy Counter Implementation (Cont.)

```
(Cont.)
24  // update: usually, just grab local lock and update local amount
25  //           once local count has risen by 'threshold', grab global
26  //           lock and transfer local values to it
27  void update(counter_t *c, int threadID, int amt) {
28      pthread_mutex_lock(&c->llock[threadID]);
29      c->local[threadID] += amt;          // assumes amt > 0
30      if (c->local[threadID] >= c->threshold) { // transfer to global
31          pthread_mutex_lock(&c->glock);
32          c->global += c->local[threadID];
33          pthread_mutex_unlock(&c->glock);
34          c->local[threadID] = 0;
35      }
36      pthread_mutex_unlock(&c->llock[threadID]);
37  }
38
39  // get: just return global amount (which may not be perfect)
40  int get(counter_t *c) {
41      pthread_mutex_lock(&c->glock);
42      int val = c->global;
43      pthread_mutex_unlock(&c->glock);
44      return val;    // only approximate!
45  }
```

Concurrent Linked Lists

```
1      // basic node structure
2      typedef struct __node_t {
3          int key;
4          struct __node_t *next;
5      } node_t;
6
7      // basic list structure (one used per list)
8      typedef struct __list_t {
9          node_t *head;
10         pthread_mutex_t lock;
11     } list_t;
12
13     void List_Init(list_t *L) {
14         L->head = NULL;
15         pthread_mutex_init(&L->lock, NULL);
16     }
```

Concurrent Linked Lists (Cont.)

```
18     int List_Insert(list_t *L, int key) {
19         pthread_mutex_lock(&L->lock);
20         node_t *new = malloc(sizeof(node_t));
21         if (new == NULL) {
22             perror("malloc");
23             pthread_mutex_unlock(&L->lock);
24             return -1; // fail
25         }
26         new->key = key;
27         new->next = L->head;
28         L->head = new;
29         pthread_mutex_unlock(&L->lock);
30         return 0; // success
31     }
32
33     int List_Lookup(list_t *L, int key) {
34         pthread_mutex_lock(&L->lock);
35         node_t *curr = L->head;
36         while (curr) {
37             if (curr->key == key) {
38                 pthread_mutex_unlock(&L->lock);
39                 return 0; // success
40             }
41             curr = curr->next;
42         }
43         pthread_mutex_unlock(&L->lock);
44         return -1; // failure
45     }
```

Concurrent Linked Lists (Cont.)

- The code **acquires** a lock in the insert routine upon entry.
- The code **releases** the lock upon exit.
 - If `malloc()` happens to *fail*, the code must also release the lock before failing the insert.
 - This kind of exceptional control flow has been shown to be **quite error prone**.
 - **Solution:** The lock and release *only surround* the actual critical section in the insert code

Concurrent Linked List: Rewritten

```
1      void List_Init(list_t *L) {
2          L->head = NULL;
3          pthread_mutex_init(&L->lock, NULL);
4      }
5
6      void List_Insert(list_t *L, int key) {
7          // synchronization not needed
8          node_t *new = malloc(sizeof(node_t));
9          if (new == NULL) {
10              perror("malloc");
11              return;
12          }
13          new->key = key;
14
15          // just lock critical section
16          pthread_mutex_lock(&L->lock);
17          new->next = L->head;
18          L->head = new;
19          pthread_mutex_unlock(&L->lock);
20      }
21
```

Concurrent Linked List: Rewritten (Cont.)

```
(Cont.)
22     int List_Lookup(list_t *L, int key) {
23         int rv = -1;
24         pthread_mutex_lock(&L->lock);
25         node_t *curr = L->head;
26         while (curr) {
27             if (curr->key == key) {
28                 rv = 0;
29                 break;
30             }
31             curr = curr->next;
32         }
33         pthread_mutex_unlock(&L->lock);
34         return rv; // now both success and failure
35     }
```

Scaling Linked List

- Hand-over-hand locking (lock coupling)
 - Add **a lock per node** of the list instead of having a single lock for the entire list.
 - When traversing the list,
 - First grabs the next node's lock.
 - And then releases the current node's lock.
 - **Enable a high degree of concurrency in list operations.**
 - However, in practice, the overheads of acquiring and releasing locks for each node of a list traversal is ***prohibitive***.

Michael and Scott Concurrent Queues

- There are two locks.
 - One for the **head** of the queue.
 - One for the **tail**.
 - The goal of these two locks is to **enable concurrency of *enqueue* and *dequeue* operations.**
- Add a dummy node
 - Allocated in the queue initialization code
 - **Enable the separation of head and tail operations**

Concurrent Queues (Cont.)

```
1     typedef struct __node_t {
2         int value;
3         struct __node_t *next;
4     } node_t;
5
6     typedef struct __queue_t {
7         node_t *head;
8         node_t *tail;
9         pthread_mutex_t headLock;
10        pthread_mutex_t tailLock;
11    } queue_t;
12
13    void Queue_Init(queue_t *q) {
14        node_t *tmp = malloc(sizeof(node_t));
15        tmp->next = NULL;
16        q->head = q->tail = tmp;
17        pthread_mutex_init(&q->headLock, NULL);
18        pthread_mutex_init(&q->tailLock, NULL);
19    }
20
```

Concurrent Queues (Cont.)

```
21     void Queue_Enqueue(queue_t *q, int value) {
22         node_t *tmp = malloc(sizeof(node_t));
23         assert(tmp != NULL);
24         tmp->value = value;
25         tmp->next = NULL;
26
27         pthread_mutex_lock(&q->tailLock);
28         q->tail->next = tmp;
29         q->tail = tmp;
30         pthread_mutex_unlock(&q->tailLock);
31     }
32
33     int Queue_Dequeue(queue_t *q, int *value) {
34         pthread_mutex_lock(&q->headLock);
35         node_t *tmp = q->head;
36         node_t *newHead = tmp->next;
37         if (newHead == NULL) {
38             pthread_mutex_unlock(&q->headLock);
39             return -1; // queue was empty
40         }
41         *value = newHead->value;
42         q->head = newHead;
43         pthread_mutex_unlock(&q->headLock);
44         free(tmp);
45         return 0;
46     }
```

Concurrent Hash Table


- Focus on a simple hash table
 - The hash table does not resize.
 - Built using the concurrent lists
 - It uses a **lock per hash bucket** each of which is represented by *a list*.

Concurrent Hash Table

```
1      #define BUCKETS (101)
2
3      typedef struct __hash_t {
4          list_t lists[BUCKETS];
5      } hash_t;
6
7      void Hash_Init(hash_t *H) {
8          int i;
9          for (i = 0; i < BUCKETS; i++) {
10              List_Init(&H->lists[i]);
11          }
12      }
13
14      int Hash_Insert(hash_t *H, int key) {
15          int bucket = key % BUCKETS;
16          return List_Insert(&H->lists[bucket], key);
17      }
18
19      int Hash_Lookup(hash_t *H, int key) {
20          int bucket = key % BUCKETS;
21          return List_Lookup(&H->lists[bucket], key);
22      }
```

Performance of Concurrent Hash Table

- From 10,000 to 50,000 concurrent updates from each of four threads.



The simple concurrent hash table **scales magnificently.**