

Operating Systems

**Condition Variable and Semaphore
(Chapter 30 ~ 31)**

Dr. Young-Woo Kwon

30. Condition Variables

Condition Variables

- There are many cases where a thread wishes to check whether a **condition** is true before continuing its execution.
- Example:
 - A parent thread might wish to check *whether a child thread has completed.*
 - This is often called a `join()`.

Condition Variables

A Parent Waiting For Its Child

```
1      void *child(void *arg) {
2          printf("child\n");
3
4          return NULL;
5      }
6
7      int main(int argc, char *argv[]) {
8          printf("parent: begin\n");
9          pthread_t c;
10         Pthread_create(&c, NULL, child, NULL); // create child
11
12         printf("parent: end\n");
13         return 0;
14     }
```

// XXX how to indicate we are done?


// XXX how to wait for child?

What we would like to see here is:

```
parent: begin
child
parent: end
```

Parent waiting for child: Spin-based Approach

```
1      volatile int done = 0;
2
3      void *child(void *arg) {
4          printf("child\n");
5          done = 1;
6          return NULL;
7      }
8
9      int main(int argc, char *argv[]) {
10         printf("parent: begin\n");
11         pthread_t c;
12         Pthread_create(&c, NULL, child, NULL); // create child
13         while (done == 0)
14             ;
15         printf("parent: end\n");
16         return 0;
17     }
```



- This is hugely inefficient as the parent spins and **wastes CPU time**.

How to wait for a condition

- Condition variable
 - Queue of threads
 - **Waiting** on the condition
 - An explicit queue that threads can put themselves on when some state of execution is not as desired.
 - **Signaling** on the condition
 - Some other thread, *when it changes its state*, can wake one of those waiting threads and allow them to continue.

Definition and Routines

- Declare condition variable

```
pthread_cond_t c;
```

- Proper initialization is required.

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);    // wait()  
pthread_cond_signal(pthread_cond_t *c);                      // signal()
```

- Operation (the POSIX calls)

- The wait() call takes a mutex as a parameter.
 - The wait() call release the lock and put the calling thread to sleep.
 - When the thread wakes up, it must re-acquire the lock.

Three variables for CV:

condition variable **c**

state variable **m**

lock **l**;

Parent waiting for Child: Use a condition variable

```
1  int done = 0;
2  pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3  pthread_cond_t c = PTHREAD_COND_INITIALIZER;
4
5  void thr_exit() {
6      pthread_mutex_lock(&m);
7      done = 1;
8      pthread_cond_signal(&c);
9      pthread_mutex_unlock(&m);
10 }
11
12 void *child(void *arg) {
13     printf("child\n");
14     thr_exit();
15     return NULL;
16 }
17
18 void thr_join() {
19     pthread_mutex_lock(&m);
20     while (done == 0)
21         pthread_cond_wait(&c, &m);
22     pthread_mutex_unlock(&m);
23 }
24
```

//3 variables

//State variable

//Signaling...

//State variable

//Waiting...

Parent waiting for Child: Use a condition variable

(cont.)

```
25     int main(int argc, char *argv[]) {  
26         printf("parent: begin\n");  
27         pthread_t p;  
28         Pthread_create(&p, NULL, child, NULL);  
29         thr_join();  
30         printf("parent: end\n");  
31         return 0;  
32     }
```

//Thread creation

//Join() can be implemented
like this

Parent waiting for Child: Use a condition variable

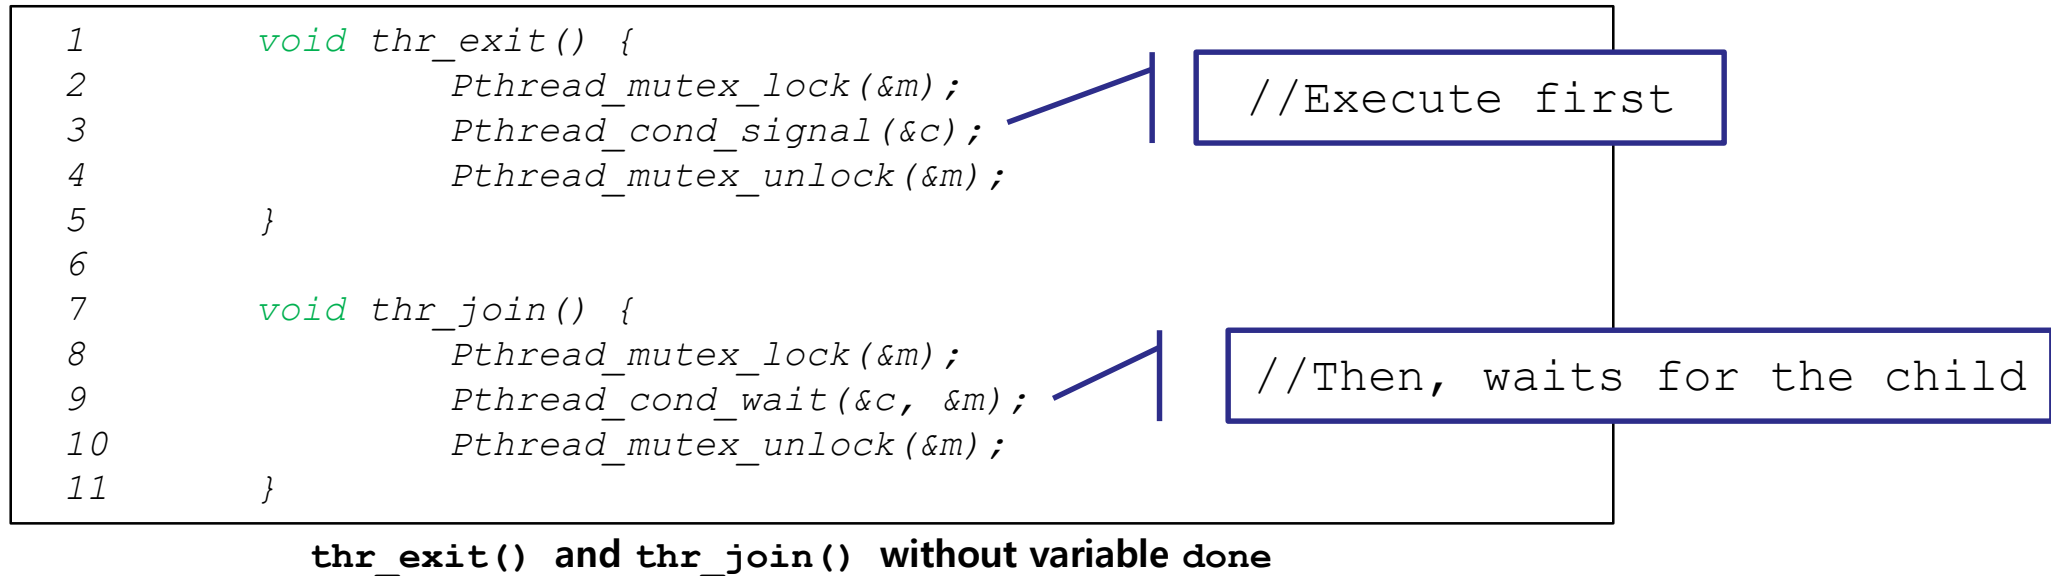
- **Parent:**

- Creates the child thread and continues running itself.
- Calls into `thr_join()` to wait for the child thread to complete.
 - Acquires the lock.
 - Checks if the child is done.
 - Puts itself to sleep by calling `wait()`.
 - Releases the lock.

- **Child:**

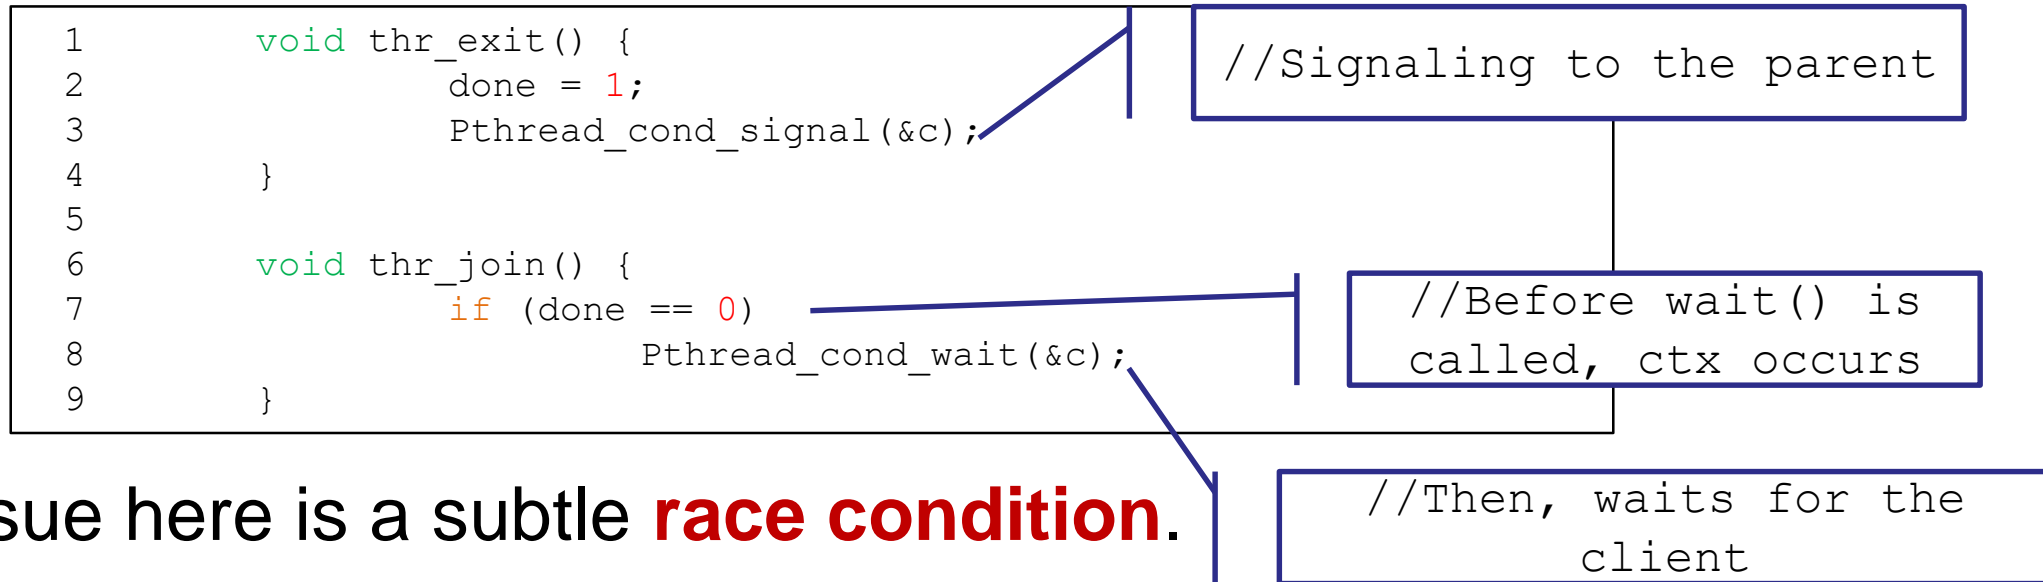
- Prints the message “child”.
- Calls `thr_exit()` to wake up the parent thread.
 - Acquire the lock.
 - Sets the state variable `done`.
 - Signals the parent thus waking it.
 - Release the lock

The importance of the state variable `done`



- Imagine the case where the *child runs immediately*.
 - The child will signal, but there is no thread asleep on the condition.
 - When the parent runs, it will call wait and be stuck.
 - No thread will ever wake it.

Another poor implementation



- The issue here is a subtle **race condition**.
 - The parent calls `thr_join()`.
 - The parent checks the value of `done`.
 - It will see that it is 0 and try to go to sleep.
 - *Just before* it calls `wait` to go to sleep, the parent is interrupted and the child runs.
 - The child changes the state variable `done` to 1 and signals.
 - But no thread is waiting and thus no thread is woken.
 - When the parent runs again, it sleeps forever.

The Producer / Consumer Problem

- **Producer**
 - **Produce** data items
 - Wish to place data items in a buffer
- **Consumer**
 - **Get** data items from the buffer and then **consume** them in some way
- Example: Multi-threaded web server
 - *A producer puts HTTP requests* in to a work queue
 - *Consumer threads take requests* out of this queue and process them

Bounded buffer

- A bounded buffer is used when you pipe the output of one program into another.
 - Example: `grep foo file.txt | wc -l`
 - The `grep` process is the producer.
 - The `wc` process is the consumer.
 - Between them is an in-kernel bounded buffer.
 - Bounded buffer is Shared resource → **Synchronized access** is required.

The Put and Get Routines (Version 1)

```
1      int buffer;  
2      int count = 0;    // initially, empty  
3  
4      void put(int value) {  
5          assert(count == 0);  
6          count = 1;  
7          buffer = value;  
8      }  
9  
10     int get() {  
11         assert(count == 1);  
12         count = 0;  
13         return buffer;  
14     }
```

- Only put data into the buffer when `count` is zero.
 - i.e., when the buffer is *empty*.
- Only get data from the buffer when `count` is one.
 - i.e., when the buffer is *full*.

Producer/Consumer Threads (Version 1)

```
1      void *producer(void *arg) {  
2          int i;  
3          int loops = (int) arg;  
4          for (i = 0; i < loops; i++) {  
5              put(i);  
6          }  
7      }  
8  
9      void *consumer(void *arg) {  
10         int i;  
11         while (1) {  
12             int tmp = get();  
13             printf("%d\n", tmp);  
14         }  
15     }
```

- **Producer** puts an integer into the shared buffer loops number of times.
- **Consumer** gets the data out of that shared buffer.

Producer/Consumer: Single CV and If Statement

- A single condition variable `cond` and associated lock `mutex`

```
1      cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              if (count == 1)                       // p2
9                  Pthread_cond_wait(&cond, &mutex); // p3
10             put(i);                                // p4
11             Pthread_cond_signal(&cond);           // p5
12             Pthread_mutex_unlock(&mutex);         // p6
13         }
14     }
15
```

//Wait until the buffer
is empty

//Signals to the consumer

Producer/Consumer: Single CV and If Statement

```
16     void *consumer(void *arg) {  
17         int i;  
18         for (i = 0; i < loops; i++) {  
19             Pthread_mutex_lock(&mutex);  
20             if (count == 0)  
21                 Pthread_cond_wait(&cond, &mutex);  
22             int tmp = get();  
23             Pthread_cond_signal(&cond);  
24             Pthread_mutex_unlock(&mutex);  
25             printf("%d\n", tmp);  
26         }  
27     }
```

//Wait until the buffer
is filled

// c3
// c4
// c5
// c6

//Signals to the producer

- p1-p3: A producer waits for the buffer to be empty.
- c1-c3: A consumer waits for the buffer to be full.
- *With just a single producer and a single consumer, the code works.*

If we have **more than** one of producer and consumer?

Thread Trace: Broken Solution (Version 1)

| | | TC1_State | | TC2_State | | TP_State | Count | Comment |
|-------------------------|----|-----------|----|-----------|----|----------|-------|--------------------|
| //Wait | c1 | Running | | Ready | | Ready | 0 | |
| | c2 | Running | | Ready | | Ready | 0 | |
| | c3 | Sleep | | Ready | | Ready | 0 | Nothing to get |
| | | Sleep | | Ready | p1 | Running | 0 | |
| | | Sleep | | Ready | p2 | Running | 0 | |
| | | Sleep | | Ready | p4 | Running | 1 | Buffer now full |
| | | Ready | | Ready | p5 | Running | 1 | awoken |
| //Switched to Tc2 by OS | | Ready | | Ready | p6 | Running | 1 | |
| | | Ready | | Ready | p1 | Running | 1 | |
| | | Ready | | Ready | p2 | Running | 1 | |
| | | Ready | | Ready | p3 | Sleep | 1 | Buffer full; sleep |
| //Woken up by TP | | Ready | c1 | Running | | Sleep | 1 | sneaks in ... |
| | | Ready | c2 | Running | | Sleep | 1 | |
| | | Ready | c4 | Running | | Sleep | 0 | ... and grabs data |
| | | Ready | c5 | Running | | Ready | 0 | awoken |
| | | Ready | c6 | Running | | Ready | 0 | |
| | c4 | Running | | Ready | | Ready | 0 | Oh oh! No data |

//Wait

Thread Trace: Broken Solution (Version 1)

- The problem arises for a simple reason:
 - After the producer woke T_{c1} , but before T_{c1} ever ran, the state of the bounded buffer *changed by* T_{c2} .
 - There is no guarantee that when the woken thread runs, the state will still be as desired → Mesa semantics.
 - **Virtually every system ever built employs *Mesa semantics*.**
 - Hoare semantics provides a stronger guarantee that the woken thread will run immediately upon being woken.

Producer/Consumer: Single CV and While

- Consumer T_{c1} wakes up and **re-checks** the state of the shared variable.
 - If the buffer is empty, the consumer simply goes back to sleep.

```
1      cond_t cond;  
2      mutex_t mutex;  
3  
4      void *producer(void *arg) {  
5          int i;  
6          for (i = 0; i < loops; i++) {  
7              Pthread_mutex_lock(&mutex);  
8              while (count == 1)  
9                  Pthread_cond_wait(&cond, &mutex);  
10             put(i);  
11             Pthread_cond_signal(&cond);  
12             Pthread_mutex_unlock(&mutex);  
13         }  
14     }  
15
```

//Check the buffer again

// p1
// p2
// p3
// p4
// p5
// p6

Producer/Consumer: Single CV and While

(Cont.)

```
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);           // c1
20             while (count == 0)                    // c2
21                 Pthread_cond_wait(&cond, &mutex); // c3
22             int tmp = get();                       // c4
23             Pthread_cond_signal(&cond);           // c5
24             Pthread_mutex_unlock(&mutex);         // c6
25             printf("%d\n", tmp);
26         }
27     }
```

//Check the buffer again

- A simple rule to remember with condition variables is to **always use while loops**.
- However, this code still has a bug (*next page*).

Thread Trace: Broken Solution (Version 2)

| | | TC1_State | | TC2_State | | TP_State | Count | Comment |
|------------------|----|-----------|----|-----------|----|----------|-------|-------------------|
| //Wait | c1 | Running | | Ready | | Ready | 0 | |
| | c2 | Running | | Ready | | Ready | 0 | |
| | c3 | Sleep | | Ready | | Ready | 0 | Nothing to get |
| //Wait | | Sleep | c1 | Running | | Ready | 0 | |
| | | Sleep | c2 | Running | | Ready | 0 | |
| | | Sleep | c3 | Sleep | | Ready | 0 | Nothing to get |
| | | Sleep | | Sleep | p1 | Running | 0 | |
| | | Sleep | | Sleep | p2 | Running | 0 | |
| //Woken by TP | | Sleep | | Sleep | p4 | Running | 1 | |
| | | Ready | | Sleep | p5 | Running | 1 | Buffer now full |
| | | Ready | | Sleep | p6 | Running | 1 | awoken |
| | | Ready | | Sleep | p1 | Running | 1 | |
| | | Ready | | Sleep | p2 | Running | 1 | |
| //Signal to ???? | | Ready | | Sleep | p3 | Sleep | 1 | Must sleep (full) |
| | c2 | Running | | Sleep | | Sleep | 1 | Recheck condition |
| | c4 | Running | | Sleep | | Sleep | 0 | grabs data |
| | c5 | Running | | Ready | | Sleep | 0 | Oops! Woke |

Thread Trace: Broken Solution (Version 2) (Cont.)

| | State | | State | | State | Count | Comment |
|-----|---------|-----|---------|-----|-------|-------|---------|
| ... | ... | ... | ... | ... | ... | ... | (cont.) |
| c6 | Running | | Ready | | Sleep | 0 | |
| c1 | Running | | Ready | | Sleep | 0 | |
| c2 | Running | | Ready | | Sleep | 0 | |
| c3 | Sleep | | Ready | | Sleep | 0 | |
| | Sleep | c2 | Running | | Sleep | 0 | |
| | Sleep | c3 | Sleep | | Sleep | 0 | |

//Woken by TC1

Nothing to get

Everyone asleep ...

- ◆ A consumer should not wake other consumers, only producers, and vice-versa.

The Single Buffer Producer/Consumer Solution

- Use **two condition variables** and while
 - **Producer** threads wait on the condition `empty`, and signals `fill`.
 - **Consumer** threads wait on `fill` and signal `empty`.

```
1  cond_t empty, fill;
2  mutex_t mutex;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          Pthread_mutex_lock(&mutex);
8          while (count == 1)
9              Pthread_cond_wait(&empty, &mutex);
10         put(i);
11         Pthread_cond_signal(&fill);
12         Pthread_mutex_unlock(&mutex);
13     }
14 }
15
```

//Wait for empty buffer

//Signal to consumers

The single Buffer Producer/Consumer Solution

(Cont.)

```
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);
20             while (count == 0)
21                 Pthread_cond_wait(&fill, &mutex);
22             int tmp = get();
23             Pthread_cond_signal(&empty);
24             Pthread_mutex_unlock(&mutex);
25             printf("%d\n", tmp);
26         }
27     }
```

//Wait for full buffer

//Signal to the producer

The Final Producer/Consumer Solution

- More **concurrency** and **efficiency** → **Add more buffer slots.**
 - Allow concurrent production or consuming to take place.
 - Reduce context switches.

```
1      int buffer[MAX];
2      int fill = 0;
3      int use = 0;
4      int count = 0;
5
6      void put(int value) {
7          buffer[fill] = value;
8          fill = (fill + 1) % MAX;
9          count++;
10     }
11
12     int get() {
13         int tmp = buffer[use];
14         use = (use + 1) % MAX;
15         count--;
16         return tmp;
17     }
```

The Final Put and Get Routines

The Final Producer/Consumer Solution (Cont.)

```
1      cond_t empty, fill;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              while (count == MAX)                 // p2
9                  Pthread_cond_wait(&empty, &mutex); // p3
10             put(i);                               // p4
11             Pthread_cond_signal(&fill);           // p5
12             Pthread_mutex_unlock(&mutex);         // p6
13         }
14     }
15
```

The Final Producer/Consumer Solution (Cont.)

```
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);           // c1
20             while (count == 0)                     // c2
21                 Pthread_cond_wait(&fill, &mutex); // c3
22             int tmp = get();                       // c4
23             Pthread_cond_signal(&empty);          // c5
24             Pthread_mutex_unlock(&mutex);         // c6
25             printf("%d\n", tmp);
26         }
27     }
```

The Final Working Solution (Cont.)

- p2: **A producer** only sleeps if all buffers are currently filled.
- c2: **A consumer** only sleeps if all buffers are currently empty.

Covering Conditions

- Assume there are zero bytes free
 - Thread T_a calls `allocate(100)`.
 - Thread T_b calls `allocate(10)`.
 - Both T_a and T_b wait on the condition and go to sleep.
 - Thread T_c calls `free(50)`.

Which waiting thread should be woken up?

Covering Conditions (Cont.)

```
1 // how many bytes of the heap are free?
2 int bytesLeft = MAX_HEAP_SIZE;
3
4 // need lock and condition too
5 cond_t c;
6 mutex_t m;
7
8 void *
9 allocate(int size) {
10     Pthread_mutex_lock(&m);
11     while (bytesLeft < size)
12         Pthread_cond_wait(&c, &m);
13     void *ptr = ...;
14     bytesLeft -= size;
15     Pthread_mutex_unlock(&m);
16     return ptr;
17 }
18
19 void free(void *ptr, int size) {
20     Pthread_mutex_lock(&m);
21     bytesLeft += size;
22     Pthread_cond_signal(&c); // whom to signal??
23     Pthread_mutex_unlock(&m);
24 }
```

//Condition re-check

// get mem from heap

//Need to wake up the
other thread?

Covering Conditions (Cont.)

- Solution (Suggested by Lampson and Redell)
 - Replace `pthread_cond_signal()` with `pthread_cond_broadcast()`
 - `pthread_cond_broadcast()`
 - Wake up **all waiting threads**.
 - Cost: too many threads might be woken.
 - Threads that shouldn't be awake will simply wake up, re-check the condition, and then go back to sleep.

31. Semaphore

Semaphore: A definition

- An object with an integer value
 - We can manipulate with two routines; `sem_wait()` and `sem_post()`.
 - Initialization

```
1  #include <semaphore.h>
2  sem_t s;
3  sem_init(&s, 0, 1); // initialize s to the value 1
```

- Declare a semaphore `s` and initialize it to the value **1**
- The second argument, 0, indicates that the semaphore is **shared** ***between threads in the same process.***

Semaphore: Interact with semaphore

- `sem_wait()`

```
1  int sem_wait(sem_t *s) {  
2      //decrement the value of semaphore s by one  
3      //wait if value of semaphore s is negative  
4  }
```

//Decrement semaphore
value by 1

- If the value of the semaphore is *one* or *higher* when called `sem_wait()`, **return right away**.
- If the value of the semaphore < 0 , it will cause the caller to suspend execution waiting for a subsequent post.
- When **negative**, the value of the semaphore is **equal to the number of waiting threads**.

Semaphore: Interact with semaphore (Cont.)

- `sem_post()`

```
1  int sem_post(sem_t *s) {  
2      increment the value of semaphore s by one  
3      if there are one or more threads waiting, wake one  
4  }
```

//Increment semaphore
value by 1

- Simply **increments** the value of the semaphore.
- If there is a thread waiting to be woken, **wakes** one of them up.

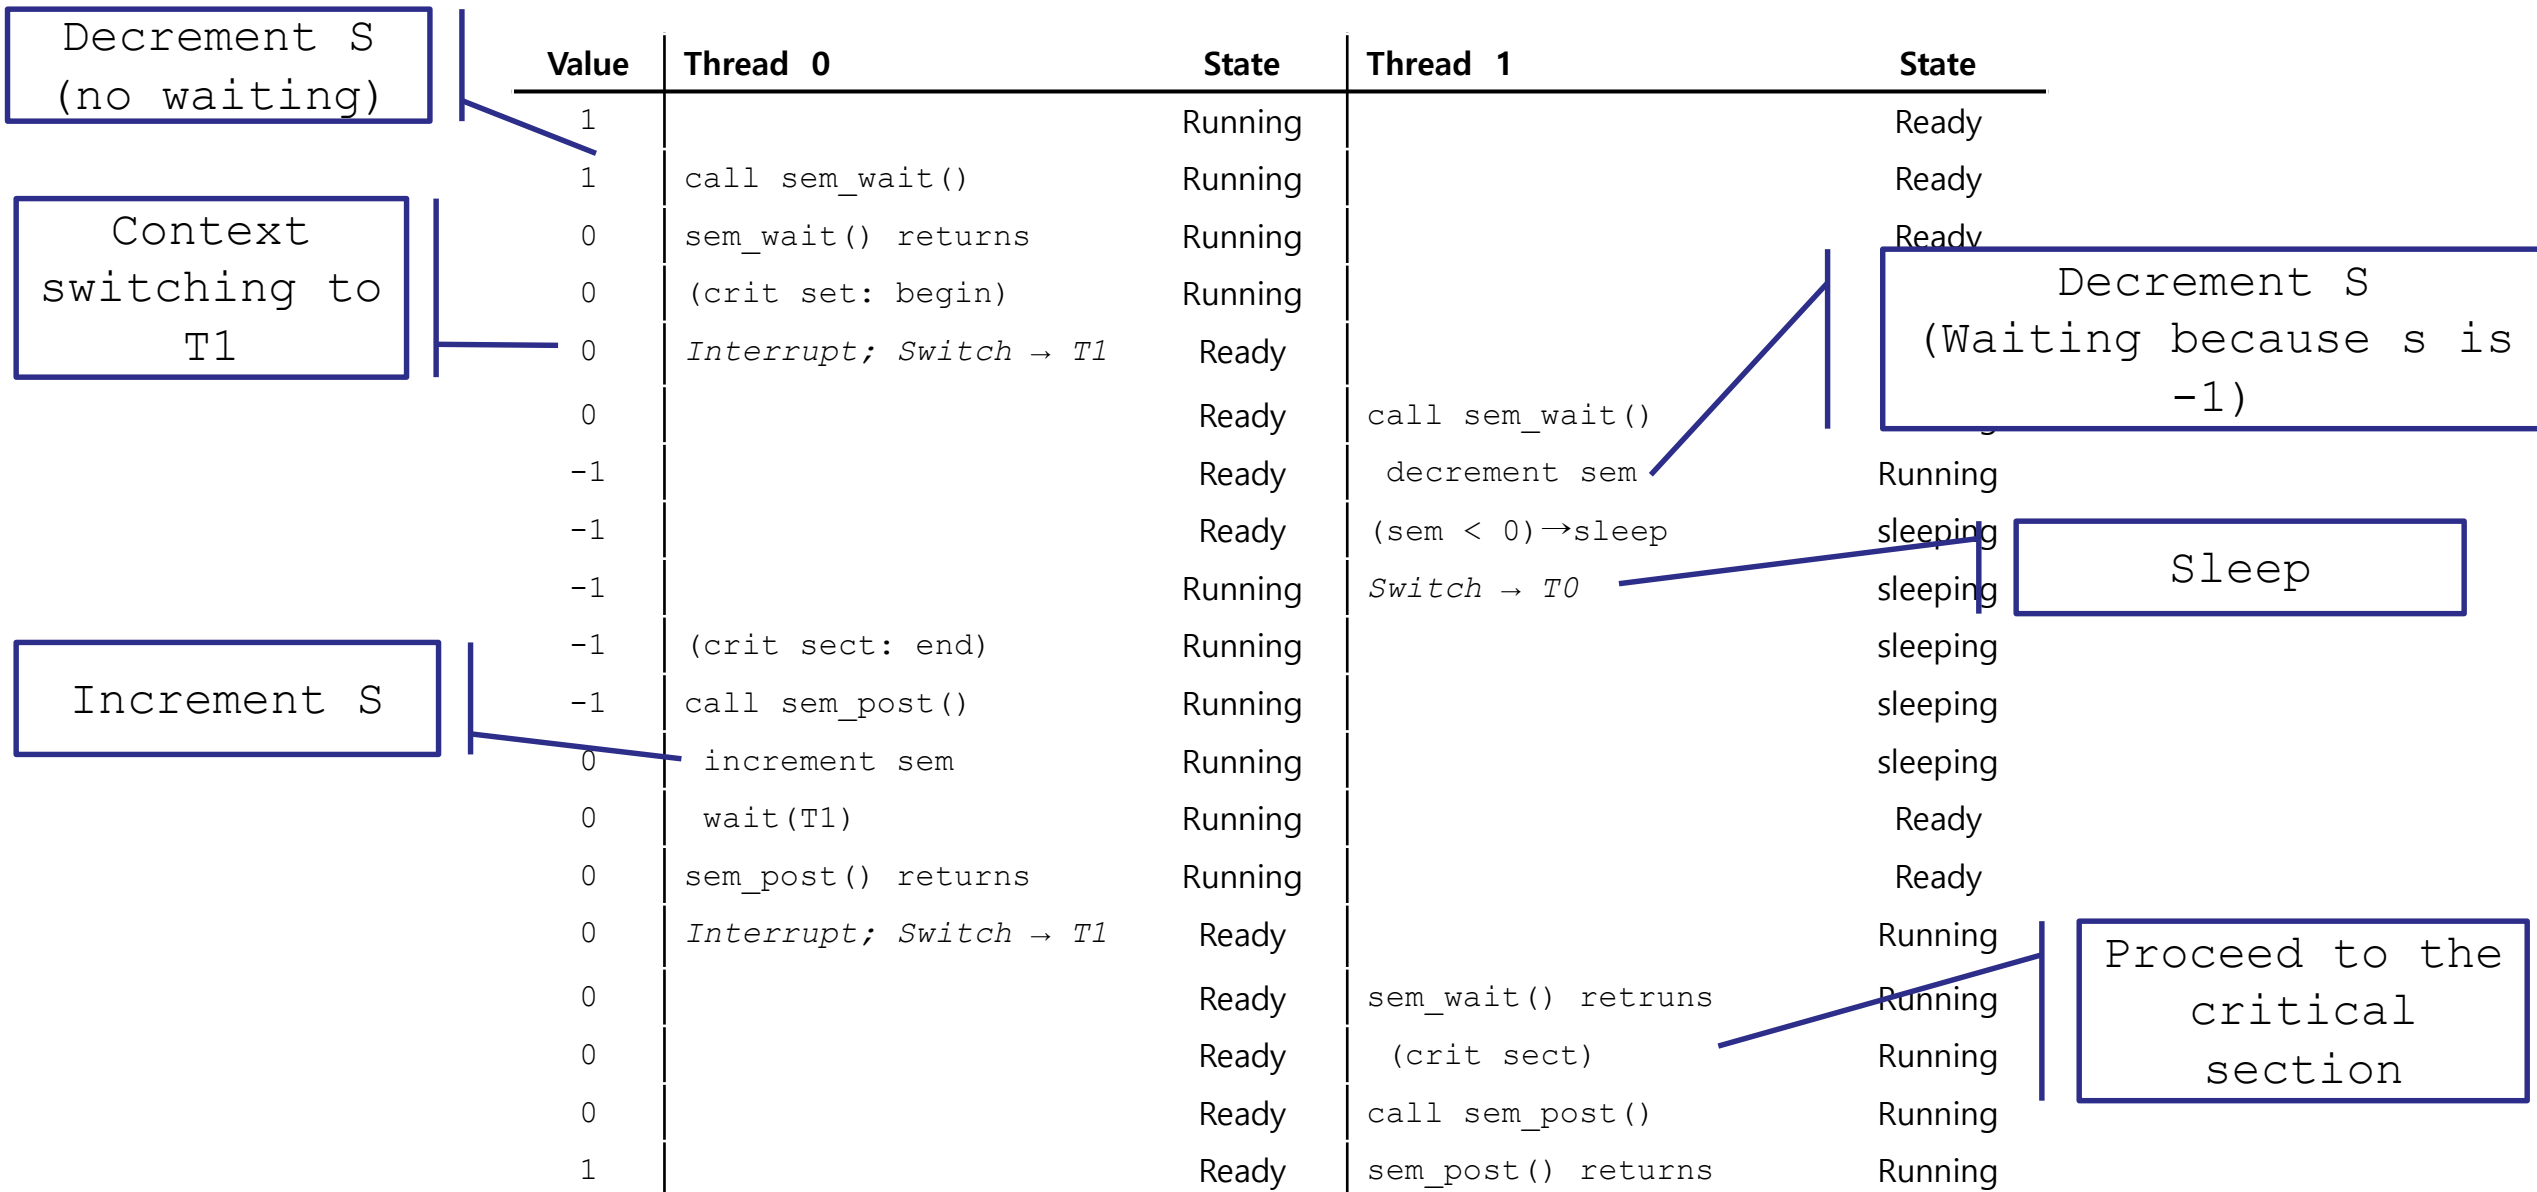
Binary Semaphores (Locks)

- What should **x** be?
 - The initial value should be **1**.

```
1  sem_t m;  
2  sem_init(&m, 0, X); // initialize semaphore to X; what should X be?  
3  
4  sem_wait(&m);  
5  //critical section here  
6  sem_post(&m);
```

| Value of Semaphore | Thread 0 | Thread 1 |
|--------------------|--------------------|----------|
| 1 | | |
| 1 | call sema_wait() | |
| 0 | sem_wait() returns | |
| 0 | (crit sect) | |
| 0 | call sem_post() | |
| 1 | sem_post() returns | |

Thread Trace: Two Threads Using A Semaphore



Semaphores as Condition Variables

```
1  sem_t s;
2
3  void *
4  child(void *arg) {
5      printf("child\n");
6      sem_post(&s); // signal here: child is done
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     sem_init(&s, 0, X); // what should X be?
13     printf("parent: begin\n");
14     pthread_t c;
15     pthread_create(c, NULL, child, NULL);
16     sem_wait(&s); // wait here for child
17     printf("parent: end\n");
18     return 0;
19 }
```

A Parent Waiting For Its Child

Semaphore variable
must be 0

parent: begin
child
parent: end

The execution result

– What should **x** be?

- The value of semaphore should be set to is **0**.

Thread Trace: Parent Waiting For Child (Case 1)

- The parent call `sem_wait()` before the child has called `sem_post()`.

| Value | Parent | State | Child | State |
|-------|--------------------------------------|----------|------------------------------------|---------|
| 0 | Create(Child) | Running | <i>(Child exists; is runnable)</i> | Ready |
| 0 | call <code>sem_wait()</code> | Running | | Ready |
| -1 | decrement sem | Running | | Ready |
| -1 | $(sem < 0) \rightarrow \text{sleep}$ | sleeping | | Ready |
| -1 | <i>Switch</i> →Child | sleeping | child runs | Running |
| -1 | | sleeping | call <code>sem_post()</code> | Running |
| 0 | | sleeping | increment sem | Running |
| 0 | | Ready | wake(Parent) | Running |
| 0 | | Ready | <code>sem_post()</code> returns | Running |
| 0 | | Ready | <i>Interrupt; Switch</i> →Parent | Ready |
| 0 | <code>sem_wait()</code> retruns | Running | | Ready |

Thread Trace: Parent Waiting For Child (Case 2)

- The child runs to completion before the parent call `sem_wait()`.

| Value | Parent | State | Child | State |
|-------|--------------------------------|---------|---------------------------------|---------|
| 0 | Create(Child) | Running | (Child exists; is runnable) | Ready |
| 0 | <i>Interrupt; switch→Child</i> | Ready | child runs | Running |
| 0 | | Ready | call sem_post() | Running |
| 1 | | Ready | increment sem | Running |
| 1 | | Ready | wake(nobody) | Running |
| 1 | | Ready | sem_post() returns | Running |
| 1 | parent runs | Running | <i>Interrupt; Switch→Parent</i> | Ready |
| 1 | call sem_wait() | Running | | Ready |
| 0 | decrement sem | Running | | Ready |
| 0 | (sem<0)→awake | Running | | Ready |
| 0 | sem_wait() retruns | Running | | Ready |

The Producer/Consumer Problem

- **Producer:** `put ()` interface
 - Wait for a buffer to become *empty* in order to put data into it.
- **Consumer:** `get ()` interface
 - Wait for a buffer to become *filled* before using it.

```
1  int buffer[MAX];
2  int fill = 0;
3  int use = 0;
4
5  void put(int value) {
6      buffer[fill] = value;    // line f1
7      fill = (fill + 1) % MAX; // line f2
8  }
9
10 int get() {
11     int tmp = buffer[use];    // line g1
12     use = (use + 1) % MAX;    // line g2
13     return tmp;
14 }
```

The Producer/Consumer Problem

```
1  sem_t empty;
2  sem_t full;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          sem_wait(&empty);           // line P1
8          put(i);                     // line P2
9          sem_post(&full);             // line P3
10     }
11 }
12
13 void *consumer(void *arg) {
14     int i, tmp = 0;
15     while (tmp != -1) {
16         sem_wait(&full);              // line C1
17         tmp = get();                  // line C2
18         sem_post(&empty);             // line C3
19         printf("%d\n", tmp);
20     }
21 }
22 ...
```

First Attempt: Adding the Full and Empty Conditions

The Producer/Consumer Problem

```
21  int main(int argc, char *argv[]) {  
22      // ...  
23      sem_init(&empty, 0, MAX);          // MAX buffers are empty to begin with...  
24      sem_init(&full, 0, 0);             // ... and 0 are full  
25      // ...  
26  }
```

First Attempt: Adding the Full and Empty Conditions (Cont.)

- Imagine that `MAX` is greater than 1 .
 - If there are multiple producers, **race condition** can happen at line *f1*.
 - It means that the old data there is overwritten.
- We've forgotten here is **mutual exclusion**.
 - The filling of a buffer and incrementing of the index into the buffer is a **critical section**.

A Solution: Adding Mutual Exclusion

```
1  sem_t empty;  
2  sem_t full;  
3  sem_t mutex;  
4  
5  void *producer(void *arg) {  
6      int i;  
7      for (i = 0; i < loops; i++) {  
8          sem_wait(&mutex);           // line p0 (NEW LINE)  
9          sem_wait(&empty);           // line p1  
10         put(i);                     // line p2  
11         sem_post(&full);             // line p3  
12         sem_post(&mutex);           // line p4 (NEW LINE)  
13     }  
14 }  
15
```

Mutex using semaphore

A Solution: Adding Mutual Exclusion

```
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         sem_wait(&mutex);           // line c0 (NEW LINE)
20         sem_wait(&full);            // line c1
21         int tmp = get();             // line c2
22         sem_post(&empty);           // line c3
23         sem_post(&mutex);           // line c4 (NEW LINE)
24         printf("%d\n", tmp);
25     }
26 }
```

Mutex using semaphore

A Solution: Adding Mutual Exclusion (Cont.)

- Imagine two thread: one producer and one consumer.
 - The consumer **acquire** the `mutex` (line c0).
 - The consumer **calls** `sem_wait()` on the full semaphore (line c1).
 - The consumer is **blocked** and **yield** the CPU.
 - The consumer still holds the mutex!
 - The producer **calls** `sem_wait()` on the binary `mutex` semaphore (line p0).
 - The producer is now **stuck** waiting too.
- a classic deadlock.

A Working Solution

```
1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&empty);           // line p1
9          sem_wait(&mutex);           // line p1.5 (MOVED MUTEX HERE...)
10         put(i);                     // line p2
11         sem_post(&mutex);           // line p2.5 (... AND HERE)
12         sem_post(&full);            // line p3
13     }
14 }
15
```

Change the order

Adding Mutual Exclusion (Correctly)

A Working Solution

```
16 void *consumer(void *arg) {
17     int i;
18     for (i = 0; i < loops; i++) {
19         sem_wait(&full);           // line c1
20         sem_wait(&mutex);          // line c1.5 (MOVED MUTEX HERE...)
21         int tmp = get();           // line c2
22         sem_post(&mutex);          // line c2.5 (... AND HERE)
23         sem_post(&empty);          // line c3
24         printf("%d\n", tmp);
25     }
26 }
27
28 int main(int argc, char *argv[]) {
29     // ...
30     sem_init(&empty, 0, MAX); // MAX buffers
31     sem_init(&full, 0, 0);    // To check if the buffer is full
32     sem_init(&mutex, 0, 1);   // mutex=1 because it is a lock
33     // ...
34 }
35 ...
```

Change the order

Adding Mutual Exclusion (Correctly)

Reader-Writer Locks

- Imagine a number of concurrent list operations, including **inserts** and simple **lookups**.
 - **insert:**
 - Change the state of the list
 - A traditional critical section makes sense.
 - **lookup:**
 - Simply *read* the data structure.
 - As long as we can guarantee that no insert is on-going, we can allow many lookups to proceed **concurrently**.

This special type of lock is known as a **reader-write lock**.

A Reader-Writer Locks

- Only **a single writer** can acquire the lock.
- Once a reader has acquired **a read lock**,
 - **More readers** will be allowed to acquire the read lock too.
 - A writer will have to wait until all readers are finished.

```
1  typedef struct _rwlock_t {
2      sem_t lock;          // binary semaphore (basic lock)
3      sem_t writelock;    // used to allow ONE writer or MANY readers
4      int readers;        // count of readers reading in critical section
5  } rwlock_t;
```

```
1. void rwlock_init(rwlock_t *rw) {
2.     rw->readers = 0;
3.     sem_init(&rw->lock, 0, 1);
4.     sem_init(&rw->writelock, 0, 1);
5. }
```

A Reader-Writer Locks

```
1. void rwlock_acquire_readlock(rwlock_t *rw) {
2.     sem_wait(&rw->lock);
3.     rw->readers++;
4.     if (rw->readers == 1)
5.         sem_wait(&rw->writelock); // first reader acquires writelock
6.     sem_post(&rw->lock);
7. }
```

Allow write lock only
for the reader

If there are multiple
readers, immediately
returns

```
1. void rwlock_release_readlock(rwlock_t *rw) {
2.     sem_wait(&rw->lock);
3.     rw->readers--;
4.     if (rw->readers == 0)
5.         sem_post(&rw->writelock); // last reader releases writelock
6.     sem_post(&rw->lock);
7. }
```

A Reader-Writer Locks (Cont.)

```
1. void rwlock_acquire_writelock(rwlock_t *rw) {  
2.     sem_wait(&rw->writelock);  
3. }
```


```
1. void rwlock_release_writelock(rwlock_t *rw) {  
2.     sem_post(&rw->writelock);  
3. }
```

A Reader-Writer Locks (Cont.)

- The reader-writer locks have **fairness problem**.
 - It would be relatively easy for reader to **starve writer**.
 - To avoid this, prevent readers from entering the lock once a writer is waiting

Dining Philosophers (Dijkstra)

- A classic
- 5 Philosophers, 1 bowl of spaghetti
- Philosophers (threads) think & eat ad infinitum
 - Need left & right fork to eat (!?)
- Want solution that prevents starvation & does not delay hungry philosophers unnecessarily



The Dining Philosophers (Cont.)

- Key challenge
 - There is **no deadlock**.
 - **No** philosopher **starves** and never gets to eat.
 - **Concurrency** is high.

```
while (1) {  
    think();  
    getforks();  
    eat();  
    putforks();  
}
```

Basic loop of each philosopher

```
// helper functions  
int left(int p) { return p; }  
  
int right(int p) {  
    return (p + 1) % 5;  
}
```

Helper functions (Downey's solutions)

- Philosopher p wishes to refer to the fork on their left \rightarrow call `left(p)`.
- Philosopher p wishes to refer to the fork on their right \rightarrow call `right(p)`.

The Dining Philosophers (Cont.)

- We need some **semaphore**, one for each fork: `sem_t forks[5]`.

```
1  void getforks() {
2      sem_wait(forks[left(p)]);
3      sem_wait(forks[right(p)]);
4  }
5
6  void putforks() {
7      sem_post(forks[left(p)]);
8      sem_post(forks[right(p)]);
9  }
```

The `getforks()` and `putforks()` Routines (Broken Solution)

– **Deadlock** occur!

- If each philosopher happens to **grab the fork on their left** before any philosopher can grab the fork on their right.
- Each will be stuck ***holding one fork and waiting for another, forever.***

A Solution: Breaking The Dependency

- Change **how forks are acquired**.
 - Let's assume that philosopher 4 acquire the forks in a *different order*.

```
1  void getforks() {  
2      if (p == 4) {  
3          sem_wait(forks[right(p)]);  
4          sem_wait(forks[left(p)]);  
5      } else {  
6          sem_wait(forks[left(p)]);  
7          sem_wait(forks[right(p)]);  
8      }  
9  }
```

- There is no situation where each philosopher grabs one fork and is stuck waiting for another. **The cycle of waiting is broken.**