

# Operating Systems

**Deadlocks (Chapter 32)**

Dr. Young-Woo Kwon

# Semaphore Implementation

- Q: What if the semaphore value is negative
- A: A negative number means the number of waiters

```
void __sched down(struct semaphore *sem)
{
    unsigned long flags;

    might_sleep();
    raw_spin_lock_irqsave(&sem->lock, flags);
    if (likely(sem->count > 0))
        sem->count--;
    else
        __down(sem);
    raw_spin_unlock_irqrestore(&sem->lock, flags);
}

static inline int __sched __down_common(struct semaphore *sem, long state,
                                         long timeout)
{
    int ret;

    trace_contention_begin(&sem, flags: 0);
    ret = __down_common(sem, state, timeout);
    trace_contention_end(&sem, ret);

    return ret;
}
```

```
static inline int __sched __down_common(struct semaphore *sem, long state,
                                         long timeout)
{
    struct semaphore_waiter waiter;

    list_add_tail(&waiter, &sem->wait_list);
    waiter.task = current;
    waiter.up = false;

    for (;;) {
        if (signal_pending_state(state, current))
            goto interrupted;
        if (unlikely(timeout <= 0))
            goto timed_out;
        __set_current_state(state);
        raw_spin_unlock_irq(&sem->lock);
        timeout = schedule_timeout(timeout);
        raw_spin_lock_irq(&sem->lock);
        if (waiter.up)
            return 0;
    }

timed_out:
    list_del(entry: &waiter, list);
    return -ETIME;

interrupted:
    list_del(entry: &waiter, list);
    return -EINTR;
}
```

# Semaphore Implementation

- Semaphore up()

```
static inline void __sched __up(struct semaphore *sem)
{
    struct semaphore_waiter *waiter = list_first_entry(&sem->wait_list,
                                                         struct semaphore_waiter, list);
    list_del(entry: &waiter->list);
    waiter->up = true;
    wake_up_process(tsk: waiter->task);
}
```

## **32. Common Concurrency Problems.**

---

# Bugs in Modern Applications

Application	What it does	Non-Deadlock	Deadlock
MySQL	Database Server	14	9
Apache	Web Server	13	4
Mozilla	Web Browser	41	16
OpenOffice	Office Suite	6	2
Total		74	31

# Non-Deadlock Bugs

- Make up a majority of concurrency bugs.
- Two major types of non deadlock bugs:
  - Atomicity violation
  - Order violation

# Atomicity-Violation Bugs

- The desired **serializability** among multiple memory accesses is *violated*.
  - Simple Example found in MySQL:
    - Two different threads access the field `proc_info` in the struct `thd`.

```
1  Thread1::  
2  if(thd->proc_info){  
3      ...  
4      fputs(thd->proc_info , ...);  
5      ...  
6  }  
7  
8  Thread2::  
9  thd->proc_info = NULL;
```

# Atomicity-Violation Bugs (Cont.)

- **Solution:** Simply add locks around the shared-variable references.

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
2
3  Thread1::
4  pthread_mutex_lock(&lock);
5  if(thd->proc_info){
6      ...
7      fputs(thd->proc_info , ...);
8      ...
9  }
10 pthread_mutex_unlock(&lock);
11
12 Thread2::
13 pthread_mutex_lock(&lock);
14 thd->proc_info = NULL;
15 pthread_mutex_unlock(&lock);
```



# Order-Violation Bugs

- The **desired order** between two memory accesses is flipped.
  - i.e., **A** should always be executed before **B**, but the order is not enforced during execution.
  - **Example:**
    - The code in Thread2 seems to assume that the variable `mThread` has already been *initialized* (and is not `NULL`).

```
1  Thread1::  
2  void init(){  
3      mThread = PR_CreateThread(mMain, ...);  
4  }  
5  
6  Thread2::  
7  void mMain(...){  
8      mState = mThread->State  
9  }
```

# Order-Violation Bugs (Cont.)

- **Solution:** Enforce ordering using **condition variables**



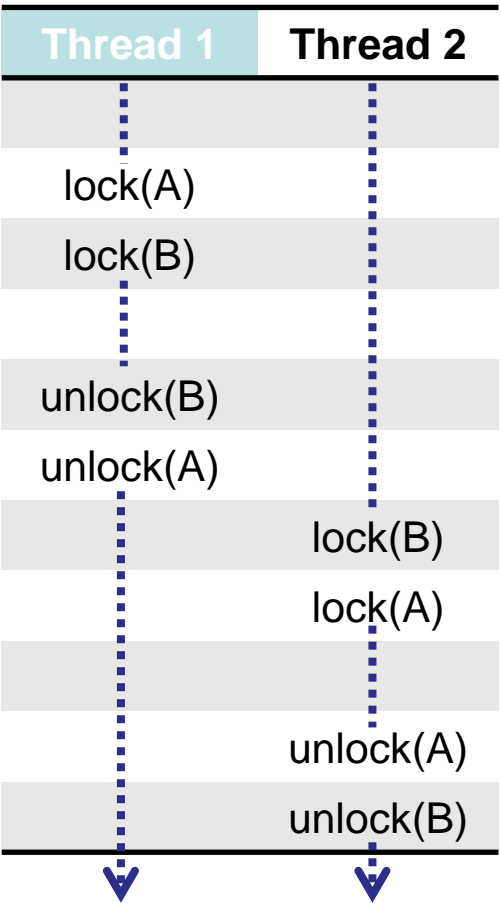
```
1  pthread_mutex_t mtLock = PTHREAD_MUTEX_INITIALIZER;
2  pthread_cond_t mtCond = PTHREAD_COND_INITIALIZER;
3  int mtInit = 0;
4
5  Thread 1::
6  void init(){
7      ...
8      mThread = PR_CreateThread(mMain,...);
9
10     // signal that the thread has been created.
11     pthread_mutex_lock(&mtLock);
12     mtInit = 1;
13     pthread_cond_signal(&mtCond);
14     pthread_mutex_unlock(&mtLock);
15     ...
16 }
17
```

# Order-Violation Bugs (Cont.)


```
18  Thread2::  
19  void mMain(...) {  
21      // wait for the thread to be initialized ...  
22      pthread_mutex_lock(&mtLock);  
23      while(mtInit == 0)  
24          pthread_cond_wait(&mtCond, &mtLock);  
25      pthread_mutex_unlock(&mtLock);  
26  
27      mState = mThread->State;  
28      ...  
29  }
```

# Deadlock Bugs


	Thread 1	Thread 2
mutex A	lock A	lock B
mutex B	lock B	lock A
	// do something	// do something
	unlock B	unlock A
	unlock A	unlock B



# Deadlock



- Simple example of circular waiting
  - Thread 1 holds lock *a*, waits on lock *b*
  - Thread 2 holds lock *b*, waits on lock *a*



# Why Do Deadlocks Occur?

- Reason 1:
  - In large code bases, **complex dependencies** arise between components.
- Reason 2:
  - Due to the nature of **encapsulation**
    - Hide details of implementations and make software easier to build in a modular way.
    - Such **modularity** *does not mesh* well with locking.

# Conditions for Deadlock

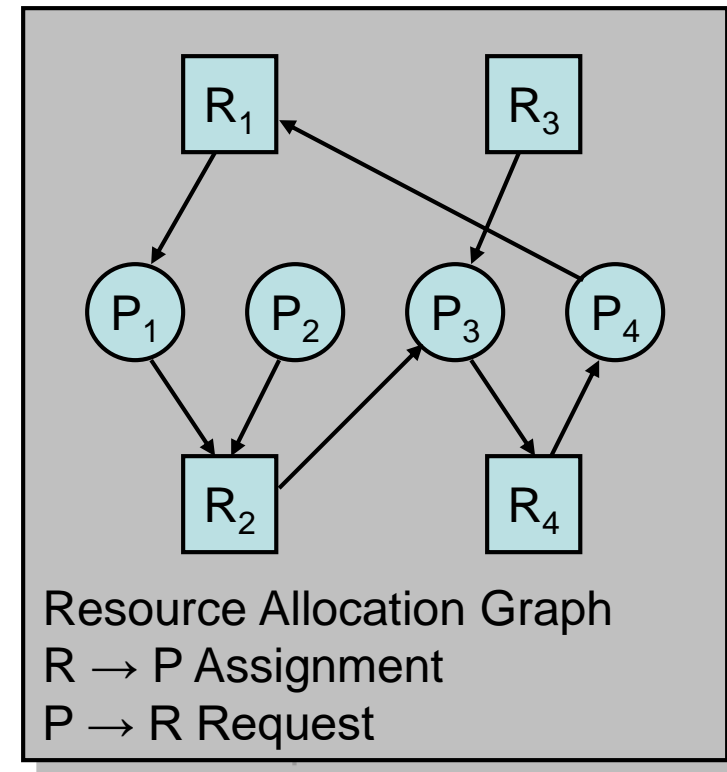
- Four conditions need to hold for a deadlock to occur.

Condition	Description
Mutual Exclusion	Threads claim exclusive control of resources that they require.
Hold-and-wait	Threads hold resources allocated to them while waiting for additional resources
No preemption	Resources cannot be forcibly removed from threads that are holding them.
Circular wait	There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain

- If any of these four conditions are not met, **deadlock cannot occur**.
- But, one more issue:
  - Buggy programming: programmer forgets to release one or more resources

# Deadlocks, more formally

- 4 necessary conditions
  - 1) Exclusive Access
  - 2) Hold and Wait
  - 3) No Preemption
  - 4) Circular Wait
  - Note that cond 1-3 represent things that are normally desirable or required
- Will look at strategies to
  - Detect & break deadlocks
  - Prevent
  - Avoid





# Detect and Recover

- **Allow deadlock** to occasionally occur and then **take some action**.
  - **Example**: if an OS froze, you would reboot it.
- Many database systems employ *deadlock detection and recovery technique*.
  - A deadlock detector **runs periodically**.
  - Building a **resource graph** and checking it for **cycles**.
  - In deadlock, the system **need to be restarted**.

# Deadlock Detection


- Idea: Look for circularity in resource allocation graph
  - Q.: How do you find out if a directed graph has a cycle?
- Can be done eagerly
  - on every resource acquisition/release, resource allocation graph is updated & tested
- or lazily
  - when all threads are blocked & deadlock is suspected, build graph & test

# Resource-Allocation Graph

- A set of vertices  $V$  and a set of edges  $E$ .
  - $V$  is partitioned into two types:
    - $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system
    - $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system
  - request edge – directed edge  $P_i \rightarrow R_j$
  - assignment edge – directed edge  $R_j \rightarrow P_i$

# Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances



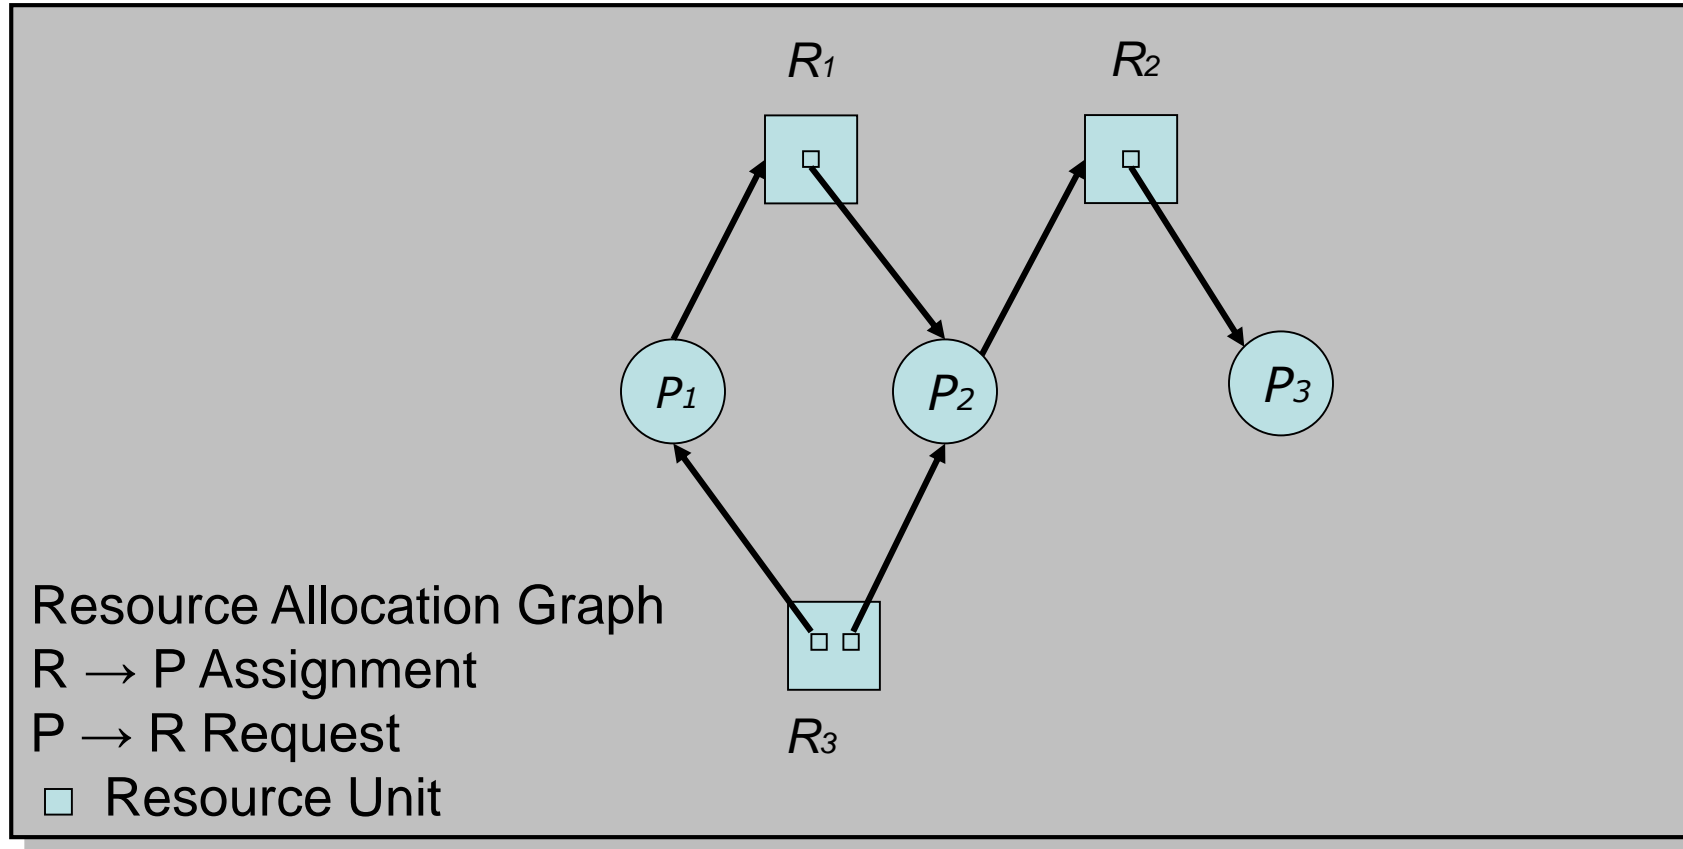
- $P_i$  requests instance of  $R_j$



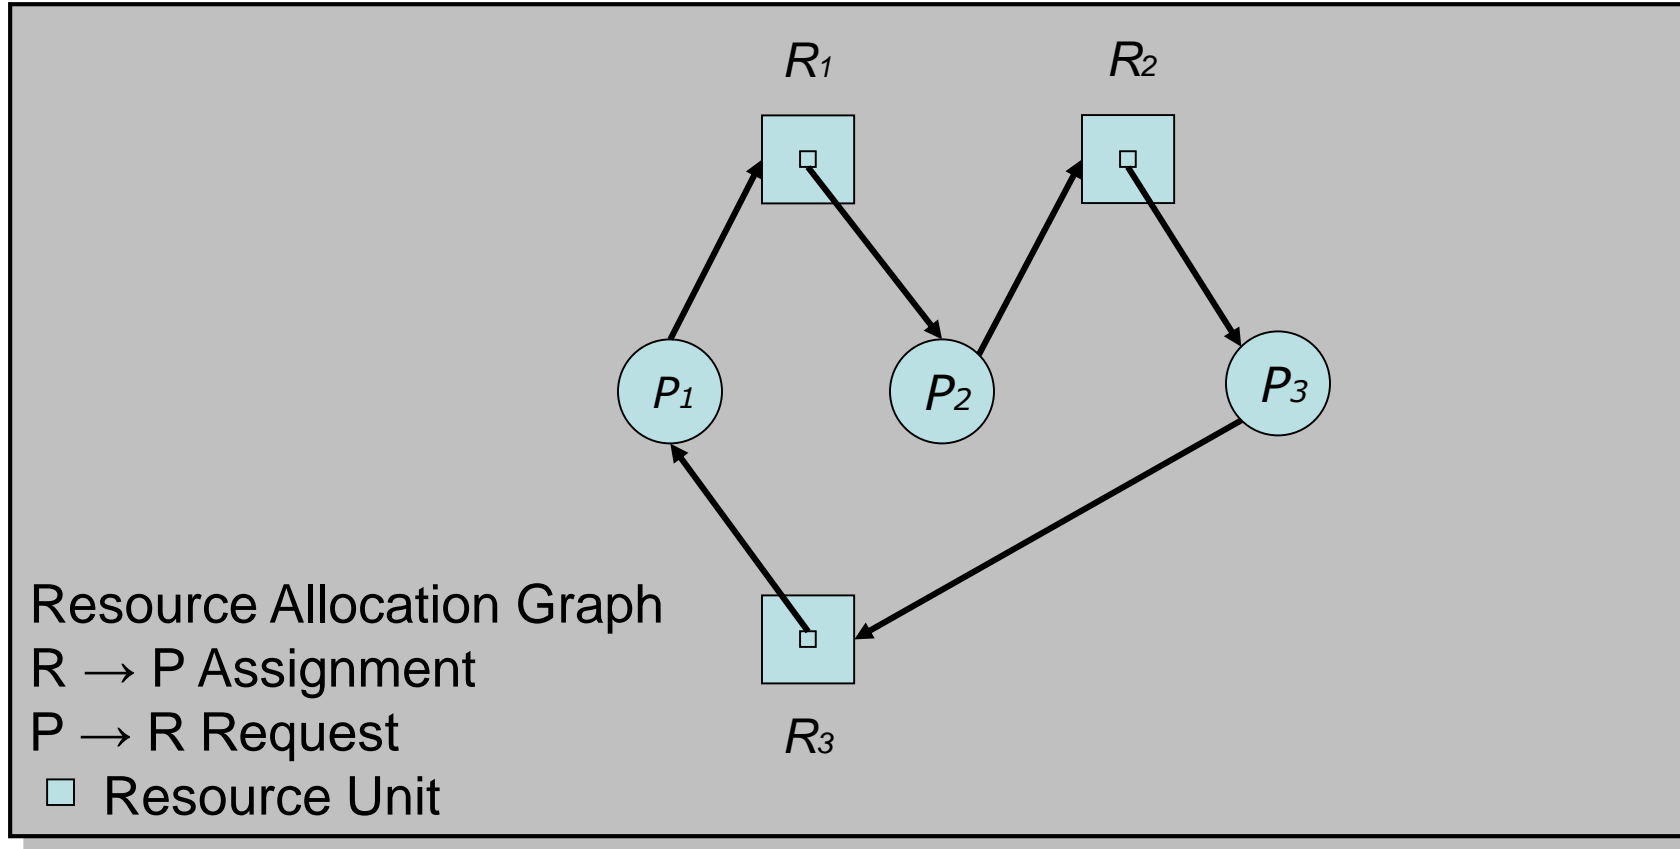
- $P_i$  is holding an instance of  $R_j$



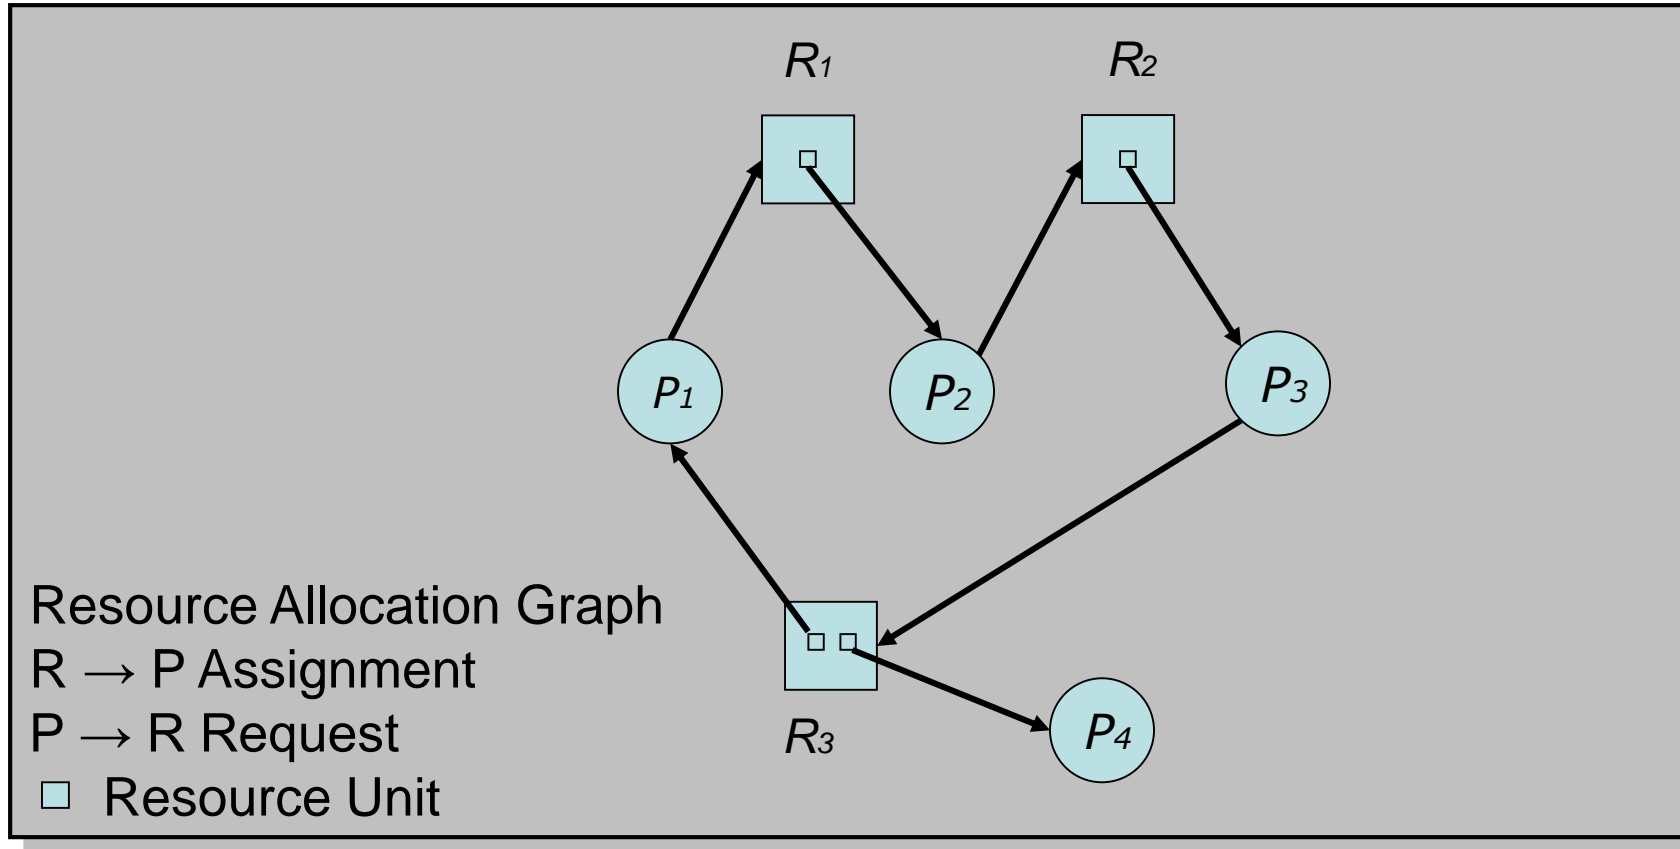
# Example of a Resource Allocation Graph



# Resource Allocation Graph with a Deadlock



# Resource Allocation Graph With A Cycle But No Deadlock



# Deadlock Detection

- Basic facts
  - If each resource has exactly one unit, deadlock iff cycle
  - If each resource has multiple units, existence of cycle may or may not mean deadlock
- Q.: What to do once deadlock is detected?



# Deadlock Recovery

Increasing Severity

- Preempt resources (if possible)
- Back processes up to a checkpoint
  - Requires checkpointing or transactions (typically expensive)
- Kill processes involved until deadlock is resolved
- Kill all processes involved
- Reboot

# Prevention – Circular Wait

- Provide **a total ordering** on lock acquisition
  - This approach requires *careful design* of global locking strategies.
- Provide **a partial ordering** in complex systems
- **Example:**
  - There are two locks in the system (L1 and L2)
  - We can prevent deadlock by always acquiring L1 before L2.

# Prevention – Hold-and-wait

- Acquire all locks **at once, atomically**.

```
1    lock(prevention);  
2    lock(L1);  
3    lock(L2);  
4    ...  
5    unlock(prevention);
```

– This code guarantees that **no untimely thread switch can occur *in the midst of*** lock acquisition.

– **Problem:**

- Require us to know when calling a routine exactly which locks must be held and to acquire them ahead of time.
- **Decrease concurrency**

# Prevention – No Preemption

- Take resource away from process
  - Difficult: how should process react?
- Virtualize resource so it can be taken away
  - Requires saving & restoring resource's state

# Prevention – No Preemption

- **Multiple lock acquisition** often gets us into trouble because when waiting for one lock **we are holding another**.
- `trylock()`
  - Used to build a *deadlock-free, ordering-robust* lock acquisition protocol.
  - Grab the lock (if it is available).
  - Or, return -1: you should try again later.

```
1  top:
2      lock(L1);
3      if( tryLock(L2) == -1 ){
4          unlock(L1);
5          goto top;
6      }
```

# Prevention – No Preemption (Cont.)

- livelock
  - Both systems are running through the code sequence *over and over again*.
  - Progress is not being made.
  - Solution:
    - Add a **random delay** before looping back and trying the entire thing over again.

# Prevention – Mutual Exclusion

- wait-free
  - Using powerful hardware instruction.
  - You can build data structures in a manner that *does not require explicit locking*.

```
1  int CompareAndSwap(int *address, int expected, int new) {
2      if(*address == expected){
3          *address = new;
4          return 1; // success
5      }
6      return 0;
7  }
```

# Prevention – Mutual Exclusion (Cont.)

- We now wanted to **atomically increment** a value by a certain amount:

```
1  void AtomicIncrement(int *value, int amount){  
2      do{  
3          int old = *value;  
4          }while( CompareAndSwap(value, old, old+amount)==0);  
5  }
```

- Repeatedly tries to update the value to *the new amount* and uses the `compare-and-swap` to do so.
- **No lock** is acquired
- **No deadlock** can arise
- **livelock** is still a possibility.



# Prevention – Mutual Exclusion (Cont.)

- **Solution:**

- Surrounding this code with a **lock acquire** and **release**.

```
1  void insert(int value){
2      node_t * n = malloc(sizeof(node_t));
3      assert( n != NULL );
4      n->value = value ;
5      lock(listlock); // begin critical section
6      n->next = head;
7      head = n;
8      unlock(listlock) ; //end critical section
9  }
```

- **wait-free manner** using the **compare-and-swap** instruction

```
1  void insert(int value) {
2      node_t *n = malloc(sizeof(node_t));
3      assert(n != NULL);
4      n->value = value;
5      do {
6          n->next = head;
7      } while (CompareAndSwap(&head, n->next, n));
8  }
```

# Deadlock Avoidance via Scheduling


- Deadlock Avoidance
  - Get the information about the locks various threads might grab during their execution.
  - schedule the threads in a way to guarantee no deadlock can occur.
- In some scenarios, **deadlock avoidance** is preferable.
- Problem: **Global knowledge** is required.

# Example of Deadlock Avoidance via Scheduling (1)

- We have two processors and four threads.
  - Lock acquisition demands of the threads:

	T1	T2	T3	T4
L1	yes	yes	no	no
L2	yes	yes	yes	no

- A smart scheduler could compute that as long as T1 and T2 are not run at the same time, **no deadlock** could ever arise.




# Example of Deadlock Avoidance via Scheduling (2)

- More contention for the same resources

	T1	T2	T3	T4
L1	yes	yes	yes	no
L2	yes	yes	yes	no

- A possible schedule that guarantees that *no deadlock* could ever occur.



- The total time to complete the jobs is lengthened considerably.

# Deadlock Avoidance via Scheduling

- Banker's algorithm
  - How much of each resource each process could possibly request
  - How much of each resource each process is currently holding
  - How much of each resource the system has available

Total resources in system:

A B C D

6 5 7 6



Processes (currently allocated resources):

A B C D

P1 1 2 2 1

P2 1 0 3 3

P3 1 2 1 0

Alloc. 3 4 6 4



Available system resources (Total – Allocated)

A B C D

3 1 1 2

Processes (maximum resources required by P):

A B C D

P1 3 3 2 2

P2 1 2 3 4

P3 1 3 5 0

Need = Max - current:

A B C D

P1 2 1 0 1

P2 0 2 0 1

P3 0 1 4 0

**Request will only be granted under below condition.**

1. If request made by process is less than equal to max need to that process.
2. If request made by process is less than equal to freely available resource in the system.

# Deadlock In The Real World

The screenshot shows the EFYtimes.com website interface. At the top, there is a red header with the site's logo and a navigation bar. Below the header, there are several links and a newsletter sign-up form. The main content area displays a database error message from Microsoft OLE DB Provider for SQL Server. The error message states that a transaction (Process ID 79) was deadlocked on lock resources and has been chosen as the deadlock victim. It advises rerunning the transaction. The error message is displayed in a white box with a black border, and the text is in a monospaced font.

**EFYtimes.com**

Free EFYtimes Newsletter  
enter your email address

IF IT'S TECHNOLOGY IT'S HERE

Monday, August 20, 2007 Archive!

<a href="#">Free Linux</a> Today's Top Choices for Free Linux on our free comparison site <a href="http://www.genieseecker.com">www.genieseecker.com</a>	<a href="#">Linux Software</a> Find Linux Software now See our Linux Software guide <a href="http://LinuxSoftware.CompareBeyond.com">LinuxSoftware.CompareBeyond.com</a>	<a href="#">Looking for Linux info?</a> Get the latest discussion threads, tutorials, and much more! <a href="http://Linux.ITtoolbox.com">Linux.ITtoolbox.com</a>	<a href="#">Small Linux PCs</a> x86 Platform Compatible Systems Excellent for Linux Applications! <a href="http://www.LogicSupply.com">www.LogicSupply.com</a>
--	--	---	--

Ads by Google

Microsoft OLE DB Provider for SQL Server error '80004005'

Transaction (Process ID 79) was deadlocked on lock resources with another process and has been chosen as the deadlock victim.  
Rerun the transaction.

/efytimes/fullnews.asp, line 66

# Deadlock in the Real World

- Most common strategy of handling deadlock
  - Test phase: fix all deadlocks detected during testing
  - Deploy phase: if deadlock happens, kill and rerun (easy!)
    - If it happens too often, or reproducibly, add deadlock detection code
- Weigh cost of **preventing** vs cost of (re-) **occurring**
- Static analysis tools detects some kinds of deadlocks before they occur
  - Idea: monitor order in which locks are taken, flag if not consistent lock order



# Deadlock vs. Starvation

- Deadlock:
  - No matter which policy the scheduler chooses, there is **no possible way for processes** to make forward progress
- Starvation:
  - There is **a possible way in which threads can make possible forward** progress, but the scheduler doesn't choose it
    - Example: strict priority scheduler will never scheduler lower priority threads as long as higher-priority thread is READY
    - Example: naïve reader/writer lock: starvation may occur by “bad luck”

# How many locks should I use?

- Could use one lock for all shared variables
  - Disadvantage: if a thread holding the lock blocks, no other thread can access *any* shared variable, even unrelated ones
- Ideally, want fine-grained locking
  - One lock only protects one (or a small set of) variables – how to pick that set?

# Multiple locks, the wrong way

```
static struct list usedlist; /* List of used blocks */
static struct list freelist; /* List of free blocks */

static struct lock alloclock; /* Protects allocations */
static struct lock freelock; /* Protects deallocations */
```

```
void *mem_alloc(...)
{
    block *b;
    lock_acquire(&alloclock);
    b = alloc_block_from_freelist();
    insert_into_usedlist(&usedlist, b);
    lock_release(&alloclock);
    return b->data;
}
```

```
void mem_free(block *b)
{
    lock_acquire(&freelock);
    list_remove(&b->elem);
    coalesce_into_freelist(&freelist, b);
    lock_release(&freelock);
}
```

Wrong: Allocating thread & deallocating thread could collide

# Multiple locks, 2<sup>nd</sup> try

```
static struct list usedlist; /* List of used blocks */
static struct list freelist; /* List of free blocks */

static struct lock usedlock; /* Protects usedlist */
static struct lock freelock; /* Protects freelist */
```

```
void *mem_alloc(...)
{
    block *b;
    lock_acquire(&freelock);
    b = alloc_block_from_freelist();
    lock_acquire(&usedlock);
    insert_into_usedlist(&usedlist, b);
    lock_release(&freelock);
    lock_release(&usedlock);
    return b->data;
}
```

```
void mem_free(block *b)
{
    lock_acquire(&usedlock);
    list_remove(&b->elem);
    lock_acquire(&freelock);
    coalesce_into_freelist(&freelist, b);
    lock_release(&usedlock);
    lock_release(&freelock);
}
```

Also wrong: deadlock!  
Always acquire multiple locks in same order -  
Or don't hold them simultaneously

# Multiple locks, correct (1)

```
static struct list usedlist; /* List of used blocks */
static struct list freelist; /* List of free blocks */

static struct lock usedlock; /* Protects usedlist */
static struct lock freelock; /* Protects freelist */
```

```
void *mem_alloc(...)
{
    block *b;
    lock_acquire(&usedlock);
    lock_acquire(&freelock);
    b = alloc_block_from_freelist();
    insert_into_usedlist(&usedlist, b);
    lock_release(&freelock);
    lock_release(&usedlock);
    return b->data;
}
```

```
void mem_free(block *b)
{
    lock_acquire(&usedlock);
    lock_acquire(&freelock);
    list_remove(&b->elem);
    coalesce_into_freelist(&freelist, b);
    lock_release(&freelock);
    lock_release(&usedlock);
}
```

Correct, but inefficient!  
Locks are always held simultaneously,  
one lock would suffice

# Multiple locks, correct (2)

```
static struct list usedlist; /* List of used blocks */
static struct list freelist; /* List of free blocks */

static struct lock usedlock; /* Protects usedlist */
static struct lock freelock; /* Protects freelist */
```

```
void *mem_alloc(...)
{
    block *b;
    lock_acquire(&freelock);
    b = alloc_block_from_freelist();
    lock_release(&freelock);
    lock_acquire(&usedlock);
    insert_into_usedlist(&usedlist, b);
    lock_release(&usedlock);
    return b->data;
}
```

```
void mem_free(block *b)
{
    lock_acquire(&usedlock);
    list_remove(&b->elem);
    lock_release(&usedlock);
    lock_acquire(&freelock);
    coalesce_into_freelist(&freelist, b);
    lock_release(&freelock);
}
```

# Conclusion

- Choosing which lock should protect which shared variable(s) is not easy – must weigh:
  - Whether all variables are always accessed together (use one lock if so)
  - Whether code inside critical section may block (if not, no throughput gain from fine-grained locking on uniprocessor)
  - Whether there is a consistency requirement if multiple variables are accessed in related sequence (must hold single lock if so)
  - Cost of multiple calls to lock/unlock (increasing parallelism advantages may be offset by those costs)